

Divvy Bike Weekly Usage Forecast

load the package:

```
library(fpp)
```

```
## Loading required package: forecast
```

```
## Registered S3 method overwritten by 'quantmod':  
##   method           from  
##   as.zoo.data.frame zoo
```

```
## Loading required package: fma
```

```
## Loading required package: expsmooth
```

```
## Loading required package: lmtest
```

```
## Loading required package: zoo
```

```
##  
## Attaching package: 'zoo'
```

```
## The following objects are masked from 'package:base':  
##  
##   as.Date, as.Date.numeric
```

```
## Loading required package: tseries
```

```
library(TSA)
```

```
## Registered S3 methods overwritten by 'TSA':  
##   method           from  
##   fitted.Arima    forecast  
##   plot.Arima     forecast
```

```
##  
## Attaching package: 'TSA'
```

```
## The following objects are masked from 'package:stats':  
##  
##   acf, arima
```

```
## The following object is masked from 'package:utils':
##
##     tar
```

```
library(tseries)
library(ggplot2)
library(forecast)
library(car)
```

```
## Loading required package: carData
```

```
library(MLmetrics)
```

```
##
## Attaching package: 'MLmetrics'
```

```
## The following object is masked from 'package:base':
##
##     Recall
```

```
library(vars)
```

```
## Loading required package: MASS
```

```
##
## Attaching package: 'MASS'
```

```
## The following objects are masked from 'package:fma':
##
##     cement, housing, petrol
```

```
## Loading required package: strucchange
```

```
## Loading required package: sandwich
```

```
## Loading required package: urca
```

Part 1: Data Processing

Import Divvy Bike data:

```
divvy_data <- read.csv('/Users/Xingkang/Desktop/divvy_data.csv')
divvy<-divvy_data$duration
```

Start aggregate the data by weekly basis

```
divvy_agg <- c()
l <- (length(divvy)-6)/7

for (i in 1:l){
  a <- (i-1)*7+1+3 # since the start day is Thursday and end day is Tuesday, we do not use the first and last three days in the dataset.
  b <- i*7+3
  c <- sum(divvy[a:b])
  #divvy_agg <- c(divvy_agg,c)
  divvy_agg[i] = c
}
```

From 2013/06/30 Sunday to 2019/12/28 Saturday; In total 339 weeks

Plot the aggregated data

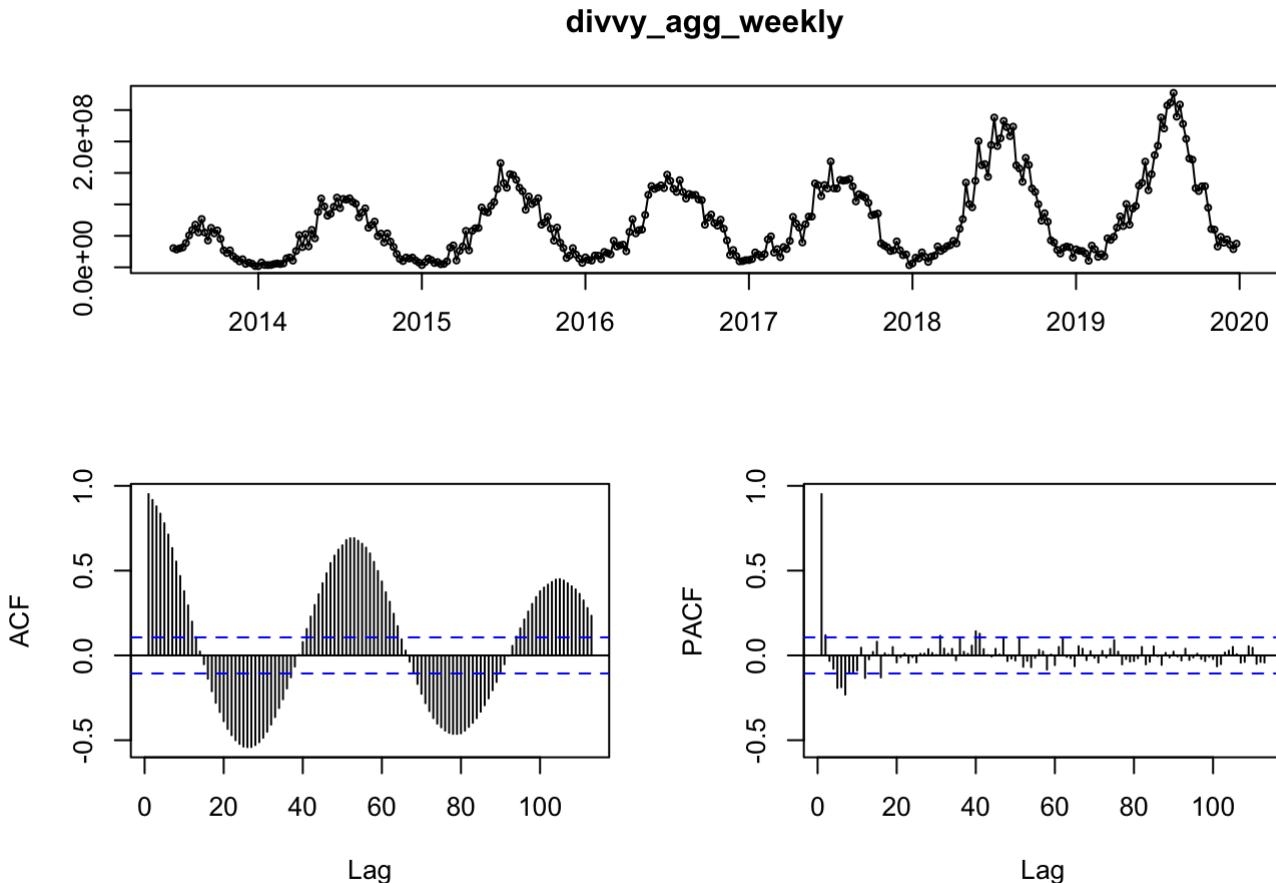
```
divvy_agg_weekly <- ts(divvy_agg, start=c(2013,26),frequency=52)
divvy_agg_weekly
```

```

## Time Series:
## Start = c(2013, 26)
## End = c(2019, 52)
## Frequency = 52
## [1] 30439829 28311019 30120450 31644021 38723962 50727234 59251259
## [8] 67795105 55483059 76746751 55198113 42622049 62657947 53968821
## [15] 58631980 45072507 26891531 22628060 26898843 17519566 13520373
## [22] 9467009 12806685 5424109 7423663 5735693 2037577 1769624
## [29] 7643736 3637395 3476235 3672188 5144109 6055938 5196131
## [36] 6182792 14298979 15627834 10649585 25880720 51193992 32085948
## [43] 52705680 33233906 59509035 46052670 88036847 109346699 96710047
## [50] 82026798 85128221 96131570 110856704 94086275 108726776 107493505
## [57] 109742674 104067226 101018086 79459035 87298888 93675136 62871412
## [64] 67141213 72684614 49552691 53533373 39515335 53858265 40954383
## [71] 32302099 21314255 13249115 9956103 15532575 13961422 15485312
## [78] 10792629 7722454 3453641 8110685 13956181 11566295 6828039
## [85] 8256944 4880240 5387374 9492880 31178979 34835921 10897601
## [92] 26153030 33127794 57879123 26904303 57871569 62041228 62518825
## [99] 95116524 88762174 87217583 97541328 103761228 124468016 165695652
## [106] 133869031 126501805 148029068 146772515 139132873 126602131 120893481
## [113] 91455600 112505923 100596514 103877142 109850897 67697792 71058907
## [120] 80322494 61370067 42409054 63183747 39034160 30538988 14847081
## [127] 18715883 30466798 20294634 13871228 6973792 15730623 12031672
## [134] 10667481 18796603 18553732 12950387 24408728 22508149 20630594
## [141] 42279014 32997876 35035773 35961257 25510119 56178451 76485787
## [148] 53982158 58946479 59973361 83307230 115197430 129145909 124578313
## [155] 126995995 130293956 125991656 147258977 137348172 124790301 119846463
## [162] 138495595 119772878 109284368 116562030 114895204 115113488 107912819
## [169] 106750137 67693242 78914234 83883029 70330831 66215186 75762207
## [176] 61257889 42793753 19507173 27408561 17851450 9505671 9845882
## [183] 11467814 11510543 13054657 23623974 20201302 16473600 20951457
## [190] 44258321 49210576 23372376 28899895 16108092 32133816 29063933
## [197] 41950159 80150389 69688401 63069970 39469918 68581655 80616026
## [204] 80633362 133322300 130653552 113127668 130802419 125380470 168238982
## [211] 125608749 125315016 138944037 137703073 138476849 140375669 128606895
## [218] 104787498 116200652 113581048 110891360 102451527 82743228 83942465
## [225] 85431815 38093795 34675408 31837263 25955102 27128811 41199915
## [232] 27202364 19263277 20423234 3313079 6185129 15316609 14075970
## [239] 23673027 16523306 8615457 16952340 18543893 33171946 25980799
## [246] 29328362 33414790 35152681 42083791 38019741 61218335 76525071
## [253] 134676910 100397490 95368995 137531594 200561134 162318031 164028059
## [260] 143841202 194417432 238100235 192821773 205257489 232747003 222801891
## [267] 208563949 223589296 162149504 157095665 135893008 174030599 162663703
## [274] 125259055 119806304 100289079 74087564 85948336 72491516 42853728
## [281] 39509845 27763304 21873709 30928069 32993725 31855264 15498967
## [288] 28600842 26271614 25799946 22180847 10426745 34445520 28418032
## [295] 16602542 20919843 17463181 46039518 43521017 48485360 63131408
## [302] 80919978 65638033 100525547 67743270 93395401 97437584 129847172
## [309] 134325875 167957018 122566812 147974818 178285109 193242332 238218202
## [316] 220713985 257302325 261689847 277309784 239483908 258729796 227934962
## [323] 204146310 172990844 171079453 125723606 121237335 128755658 128838632
## [330] 94832753 60821586 59900414 32995787 48439660 39050558 44207862
## [337] 35896430 28824347 37586842

```

```
tsdisplay(divvy_agg_weekly)
```



Split Train and Test (287 Weeks for Train; 52 Weeks for Test)

```
train_divvy <- window(divvy_agg_weekly,start=c(2013,26),end=c(2018,52))
test_divvy <- window(divvy_agg_weekly,start=c(2019,1),end=c(2019,52))
```

Import Weather Data and Split it the same way as Divvy Bike data

```
weather <- read.csv('/Users/Xingkang/Desktop/weather_data.csv')

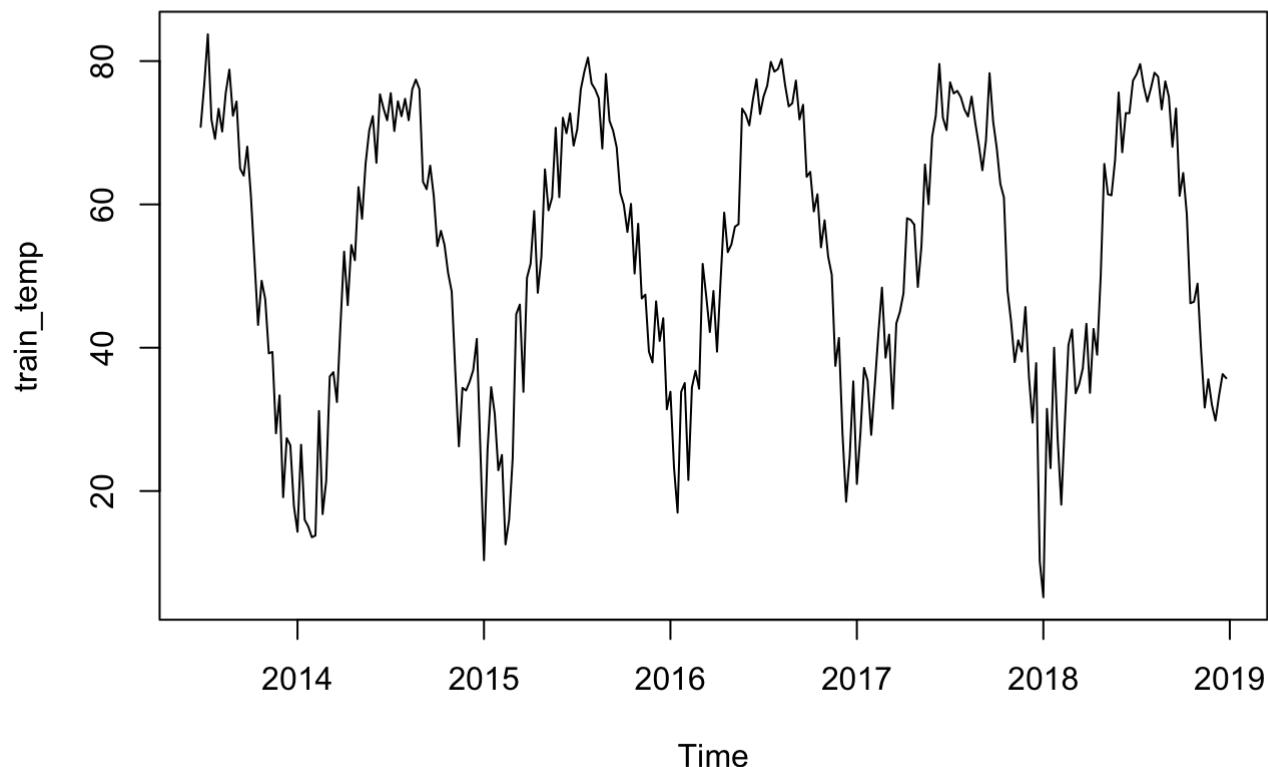
temperature <- ts(weather$avg_temp,start=c(2013,26),frequency=52)
precipitation <- ts(weather$sum_pre,start=c(2013,26),frequency=52)
snowdepth <- ts(weather$sum_snow,start=c(2013,26),frequency=52)
windspeed <- ts(weather$avg_wind,start=c(2013,26),frequency=52)

train_temp <- window(temperature,start=c(2013,26),end=c(2018,52))
train_pre <- window(precipitation,start=c(2013,26),end=c(2018,52))
train_snow <- window(snowdepth,start=c(2013,26),end=c(2018,52))
train_wind <- window(windspeed,start=c(2013,26),end=c(2018,52))

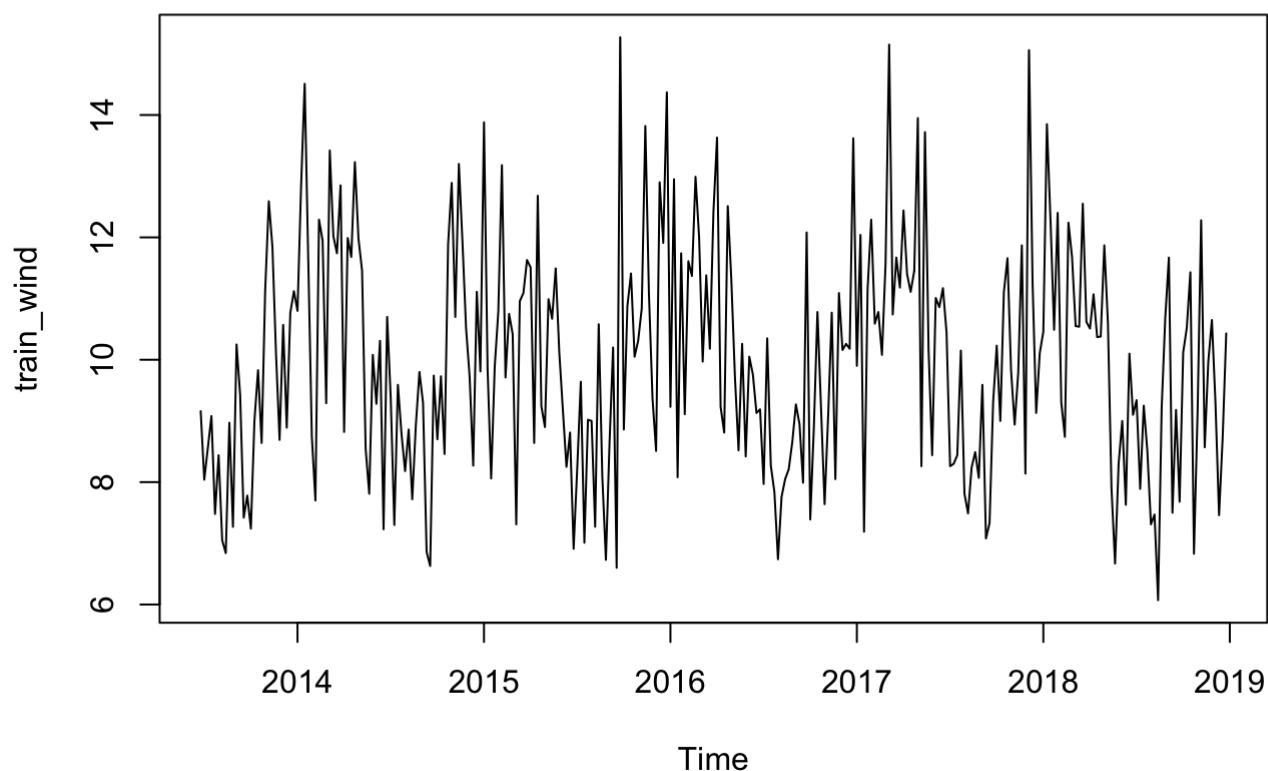
test_temp <- window(temperature,start=c(2019,1),end=c(2019,52))
```

Plot each of the 4 variables in weather dataset

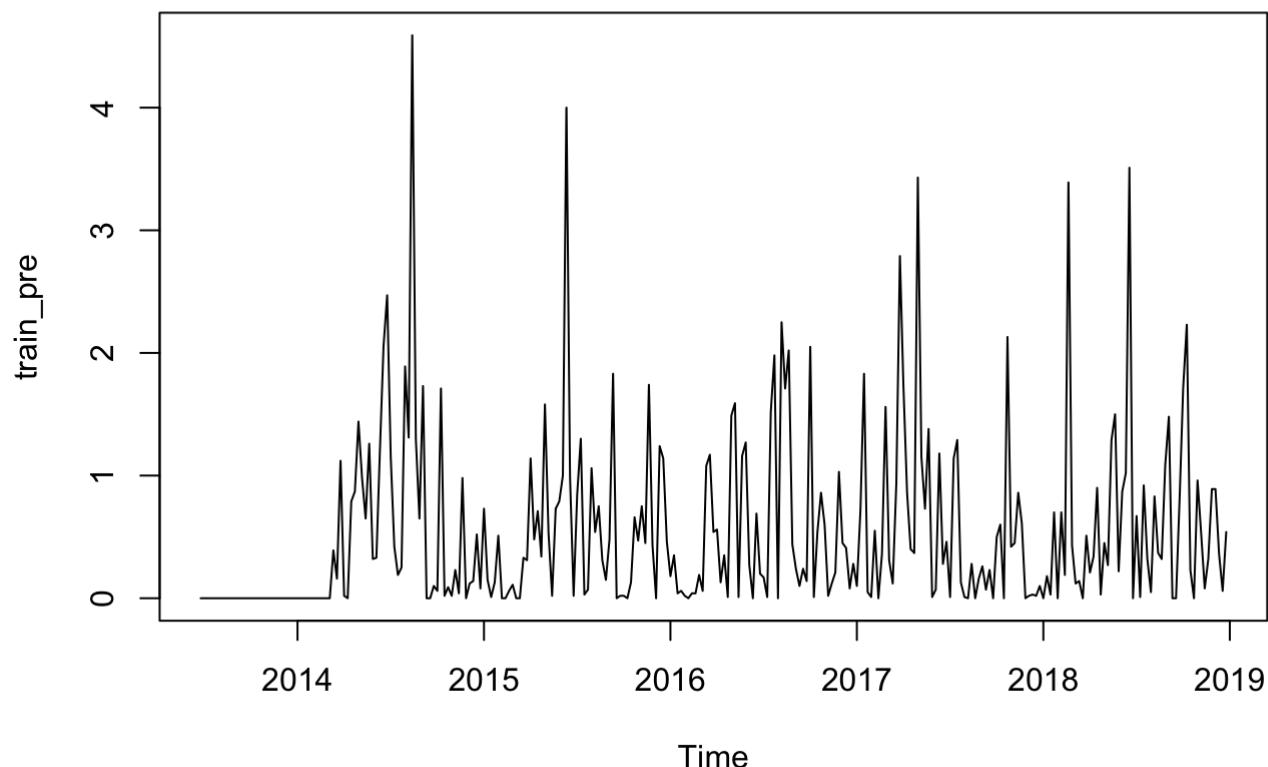
```
plot(train_temp)
```



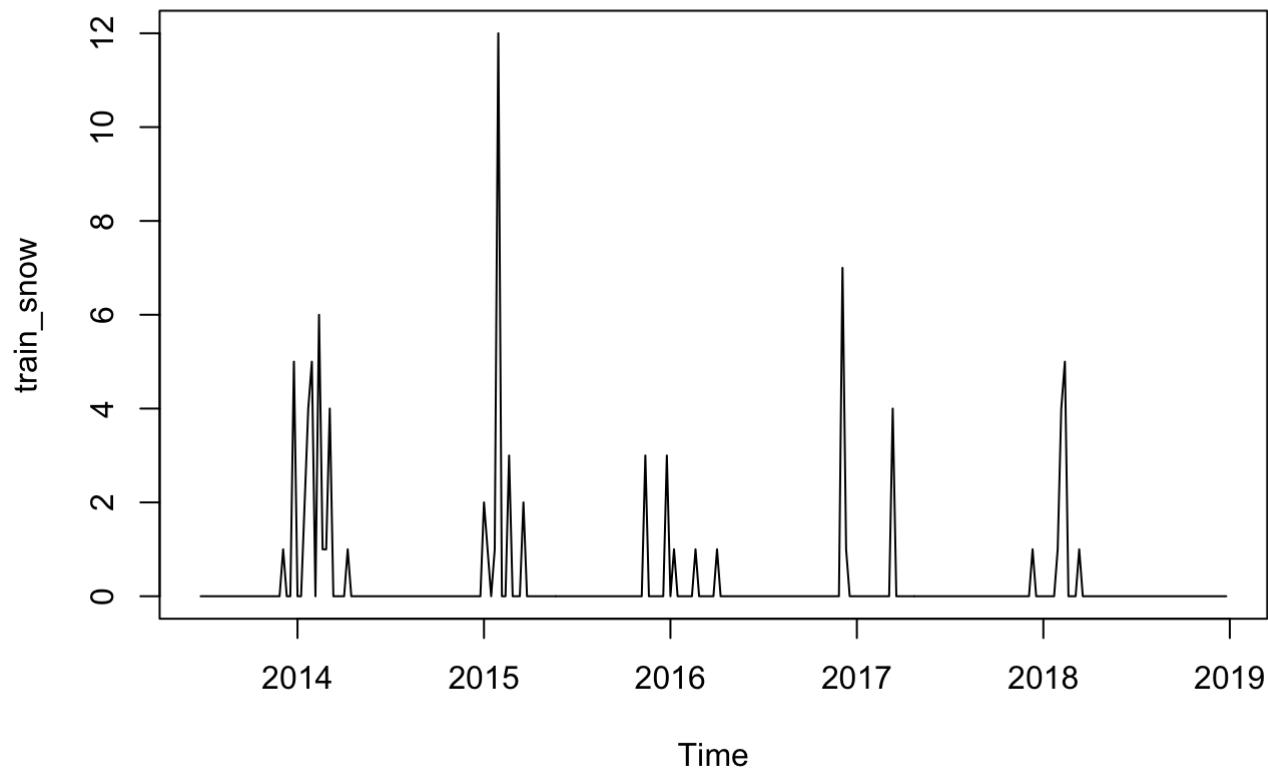
```
plot(train_wind)
```



```
plot(train_pre)
```

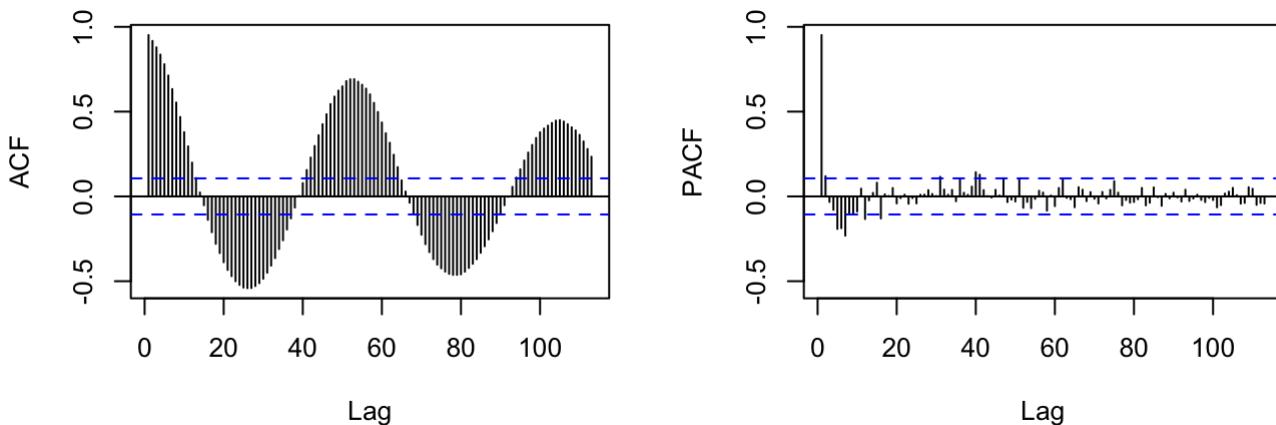
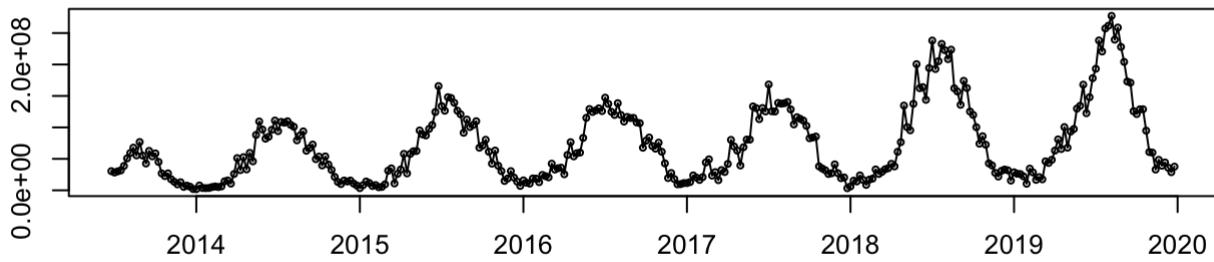


```
plot(train_snow)
```



Part 2: EDA

```
divvy_agg_weekly <- ts(divvy_agg, start=c(2013,26),frequency=52)
tsdisplay(divvy_agg_weekly)
```

divvy_agg_weekly

```
length(divvy_agg_weekly)
```

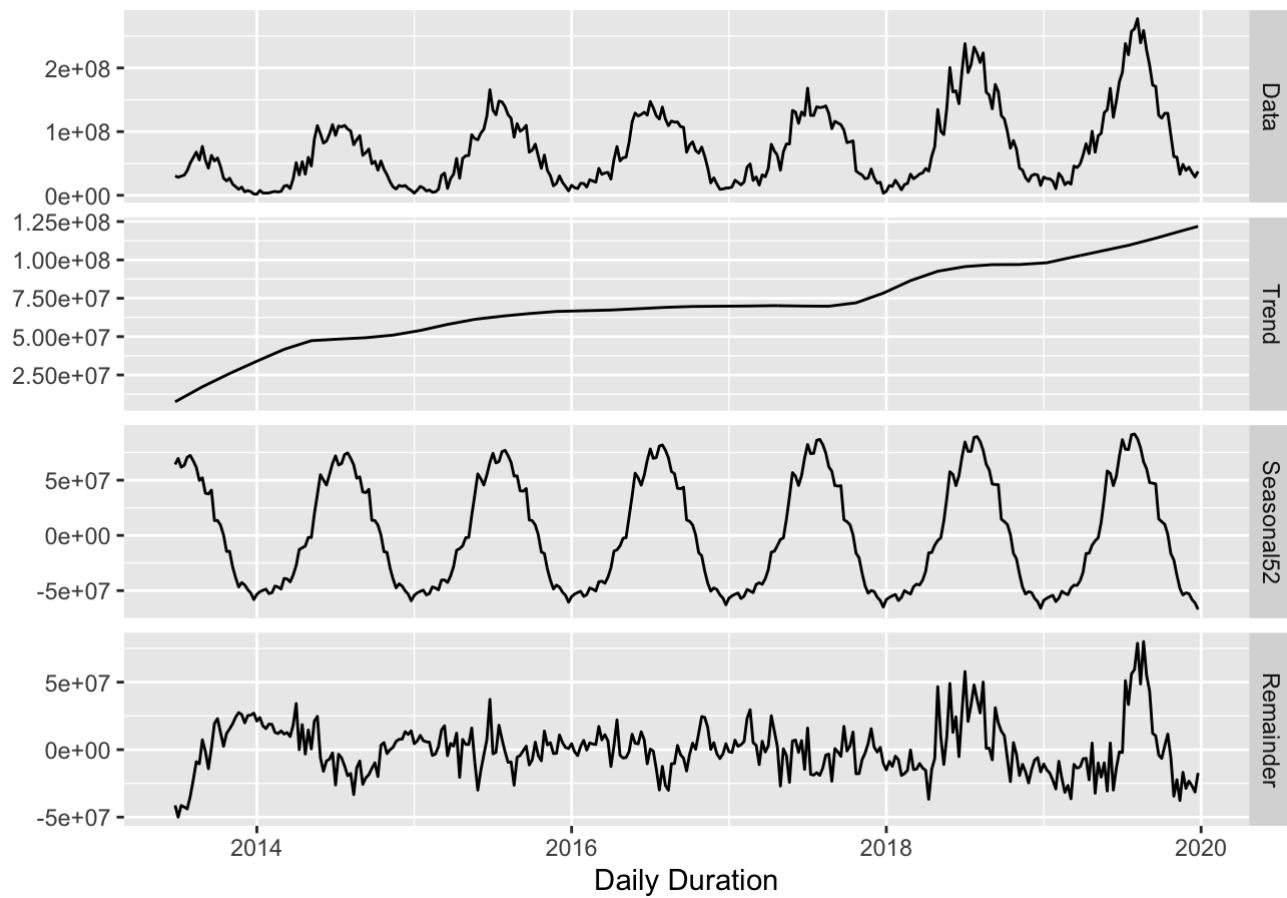
```
## [1] 339
```

```
str(divvy_agg_weekly)
```

```
## Time-Series [1:339] from 2013 to 2020: 30439829 28311019 30120450 31644021 38723962
50727234 59251259 67795105 55483059 76746751 ...
```

Time Series Decomposition Plots

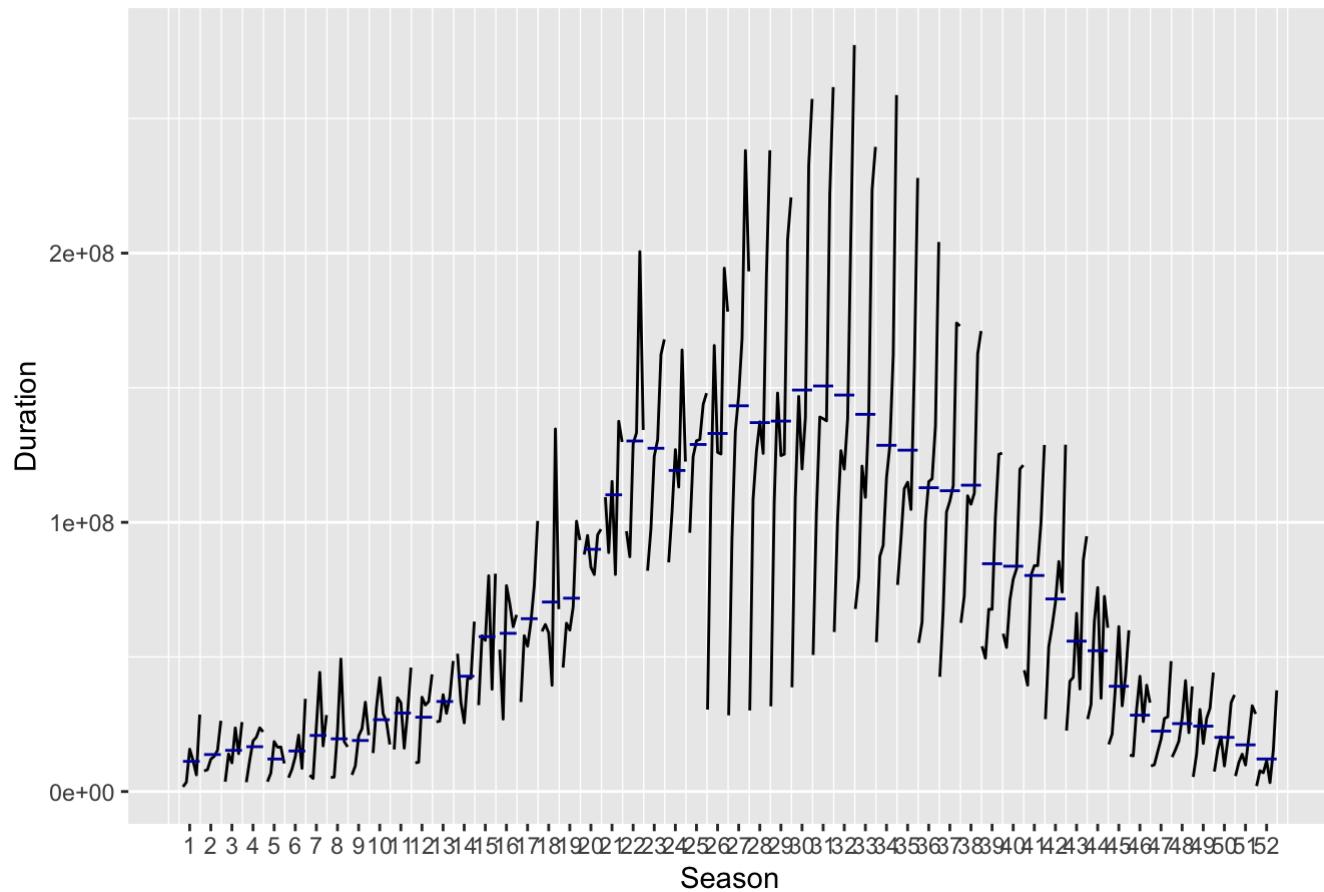
```
divvy_agg_weekly %>% mstl() %>%
  autoplot() + xlab("Daily Duration")
```



Seasonal Subseries Plots

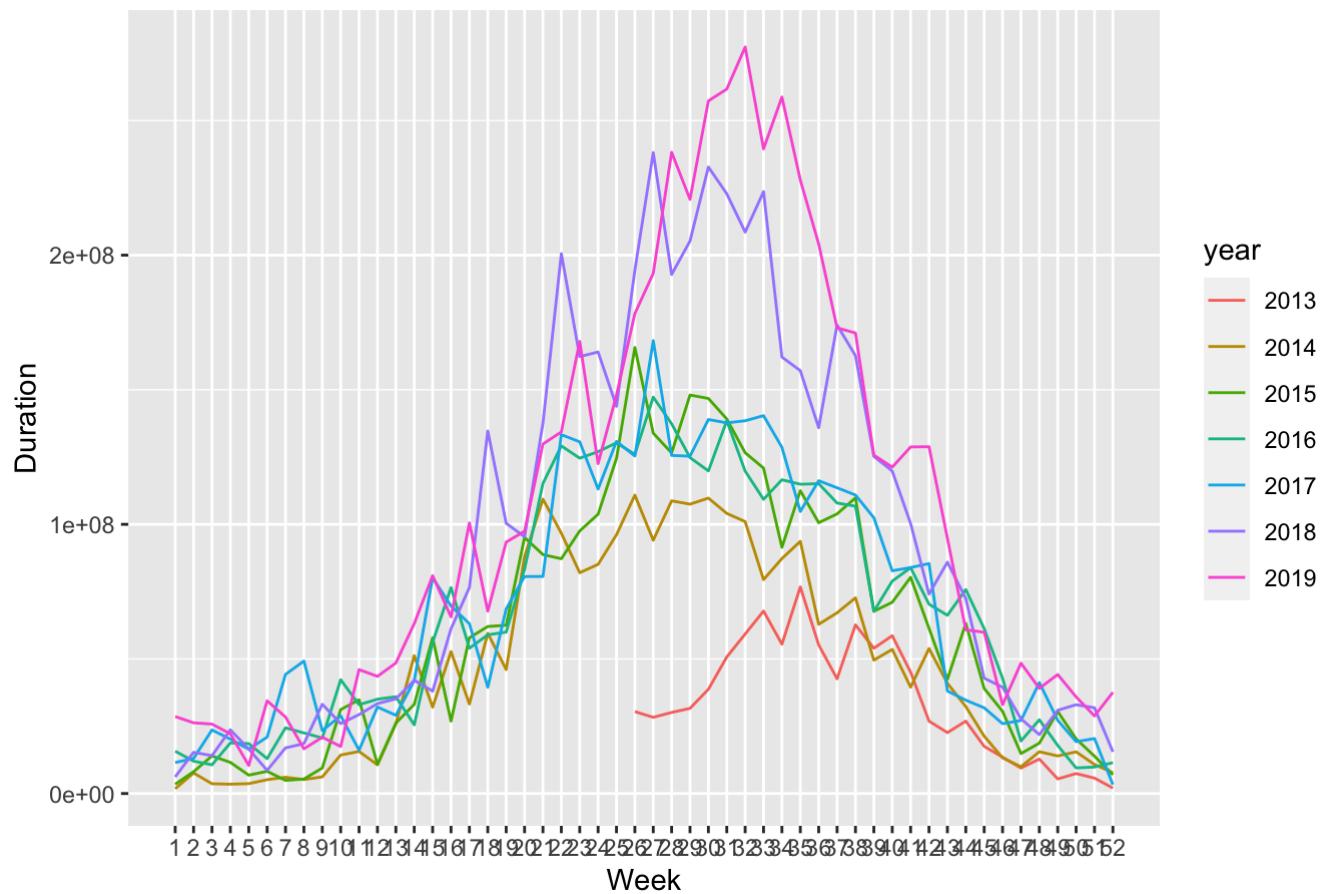
```
ggsubseriesplot(divvy_agg_weekly) +
  ylab("Duration") +
  ggtitle("Seasonal subseries plot: Divvy Weekly Duration")
```

Seasonal subseries plot: Divvy Weekly Duration



```
ggseasonplot(divvy_agg_weekly) +  
  ylab("Duration") +  
  ggtitle("Seasonal plot: Divvy Daily Duration")
```

Seasonal plot: Divvy Daily Duration

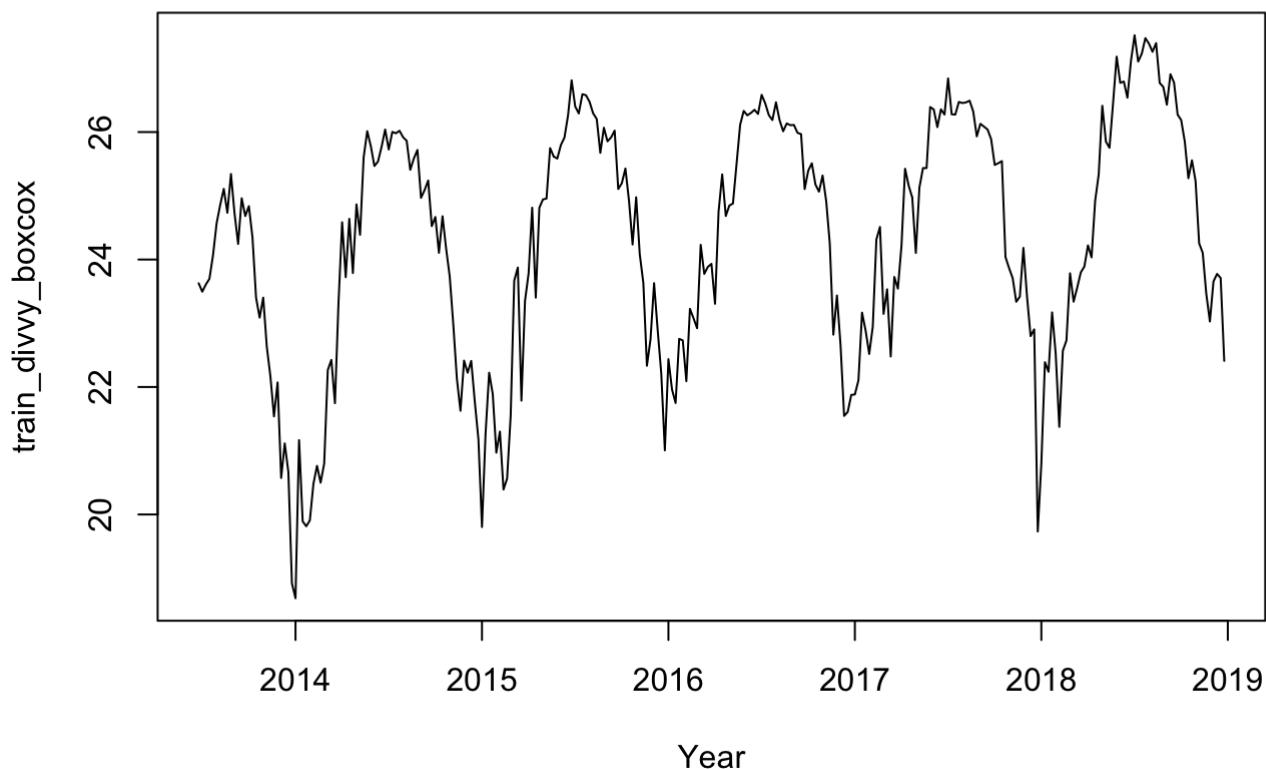


Box-Cox transformation

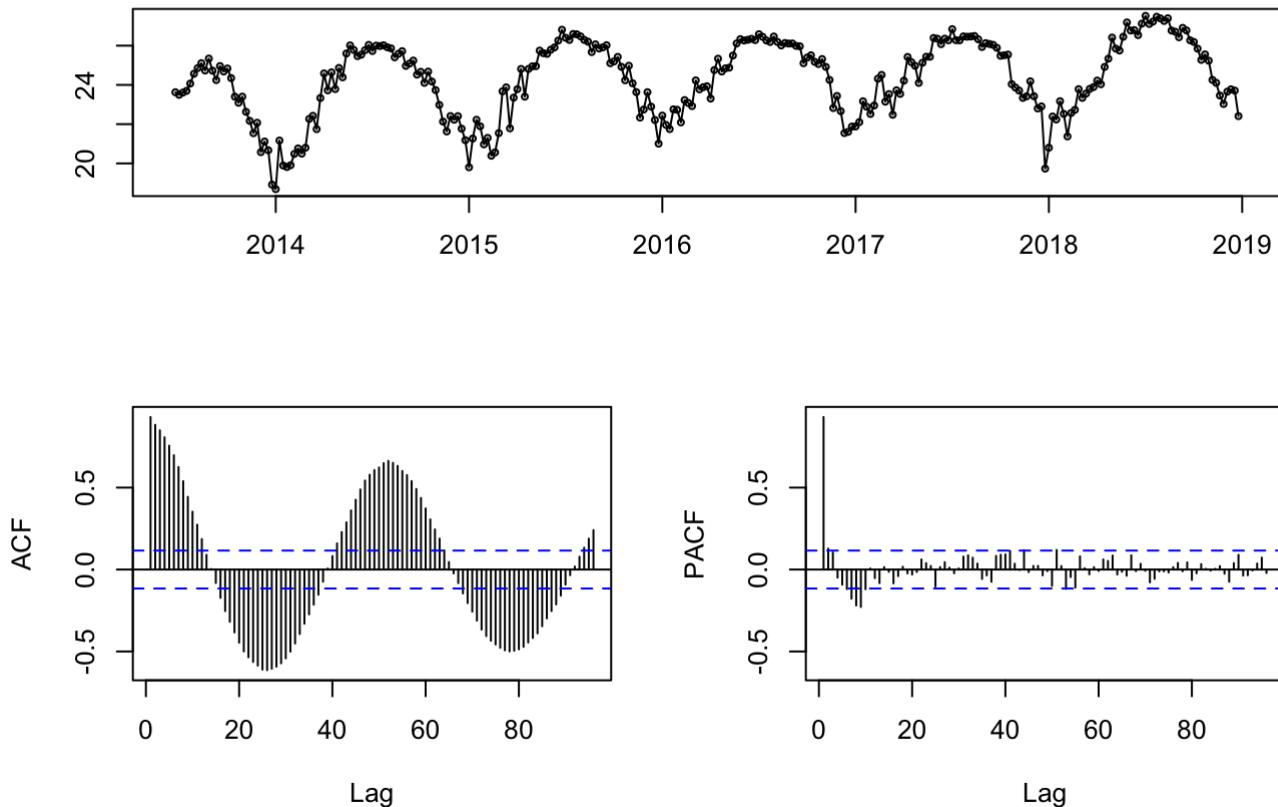
```
BoxCox.lambda(train_divvy)
```

```
## [1] 0.03489689
```

```
train_divvy_boxcox <- BoxCox(train_divvy, lambda = 0.03489689)
plot(train_divvy_boxcox, plot.type="single", col=1:2, xlab="Year")
```



```
tsdisplay(train_divvy_boxcox)
```

train_divvy_boxcox

Check if trend stationary

```
kpss.test(train_divvy_boxcox) ## small p; not trend stationary
```

```
##  
## KPSS Test for Level Stationarity  
##  
## data: train_divvy_boxcox  
## KPSS Level = 0.59309, Truncation lag parameter = 5, p-value = 0.02326
```

p-value < 0.05, the null hypothesis of trend being stationary is rejected. Therefore, we apply seasonal differencing.

```
train_divvy_boxcox_seadiff <- diff(train_divvy_boxcox, lag=52)
```

Check again if the transformed series is stationary.

```
kpss.test(train_divvy_boxcox_seadiff)
```

```
## Warning in kpss.test(train_divvy_boxcox_seadiff): p-value smaller than printed  
## p-value
```

```
##  
## KPSS Test for Level Stationarity  
##  
## data: train_divvy_boxcox_seadiff  
## KPSS Level = 0.88875, Truncation lag parameter = 4, p-value = 0.01
```

small p; trend not stationary after D=1, then we try first order differencing.

```
train_divvy_boxcox_seadiff_diff <- diff(train_divvy_boxcox_seadiff)
```

Check again if the transformed series is stationary.

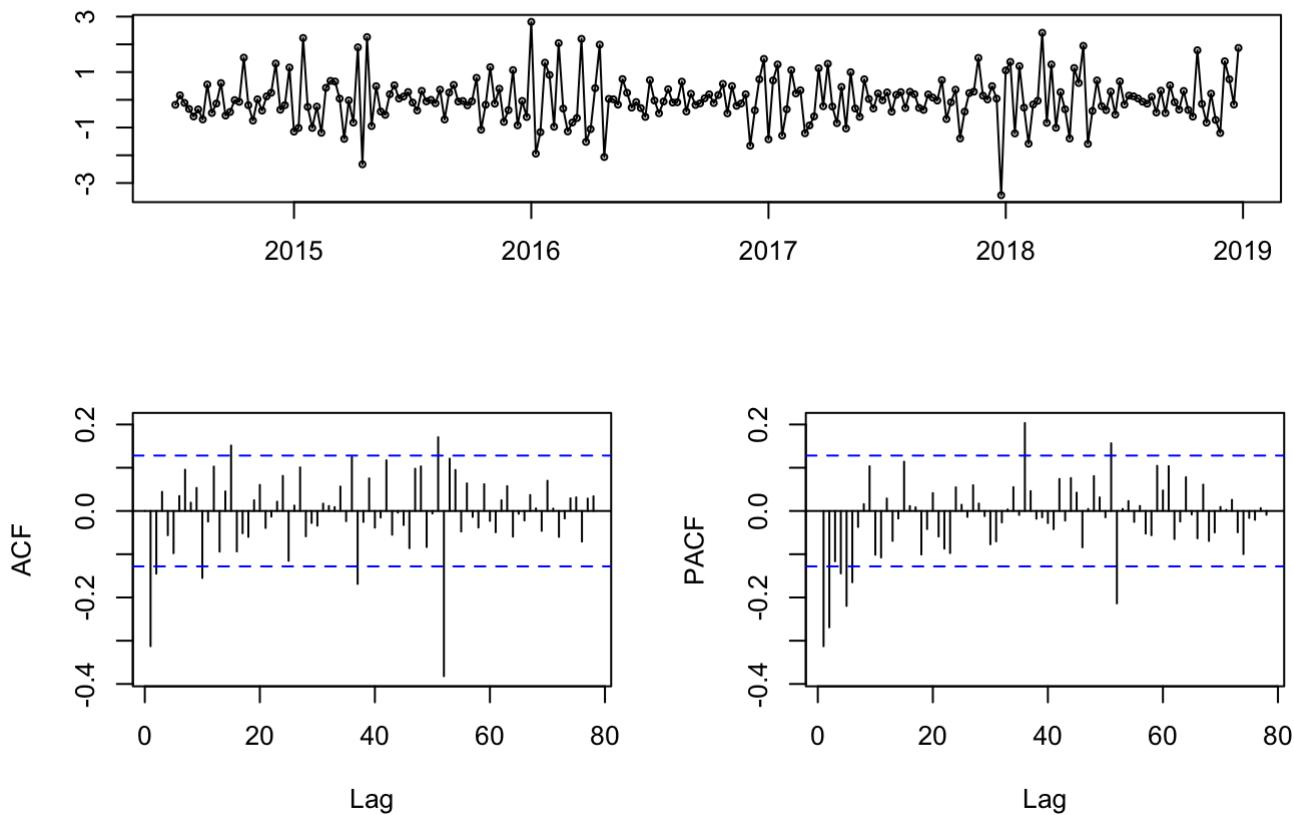
```
kpss.test(train_divvy_boxcox_seadiff_diff)
```

```
## Warning in kpss.test(train_divvy_boxcox_seadiff_diff): p-value greater than  
## printed p-value
```

```
##  
## KPSS Test for Level Stationarity  
##  
## data: train_divvy_boxcox_seadiff_diff  
## KPSS Level = 0.08886, Truncation lag parameter = 4, p-value = 0.1
```

large p; trend stationary; d=1

```
tsdisplay(train_divvy_boxcox_seadiff_diff)
```

train_divvy_boxcox_seadiff_diff

The divvy bike sereis needs seasonal differencing of lag = 52 and the first order differencing.

Check correlation between weather variables and divvy useage

```
cor(train_temp,train_divvy)
```

```
## [1] 0.8330095
```

```
cor(train_pre,train_divvy)
```

```
## [1] 0.1683076
```

```
cor(train_wind,train_divvy)
```

```
## [1] -0.4791209
```

```
cor(train_snow,train_divvy)
```

```
## [1] -0.2632397
```

The “divvy vs avg_temp” has a high positive correlation. The second highest correlation is “divvy vs avg_wind”, which has a negative correlation. However, its Pearson r is only 0.48, indicating that the correlation is not significant. Therefore, we would only use temperature as additional variable going forward

Part 3: Seasonal Naive

Build Model and get prediction

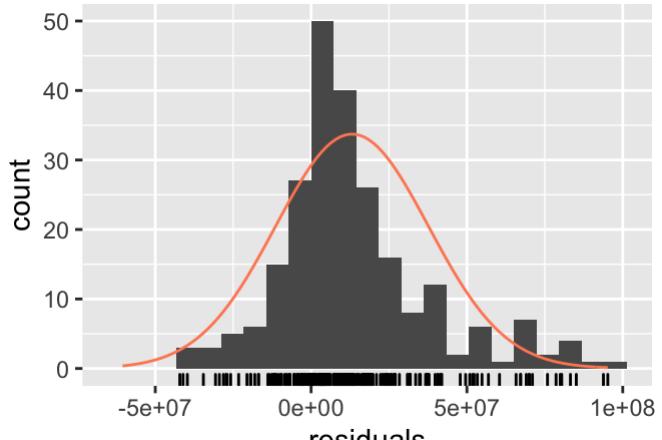
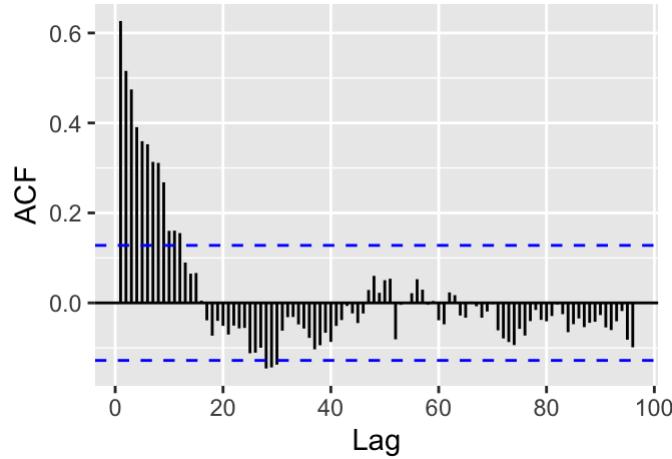
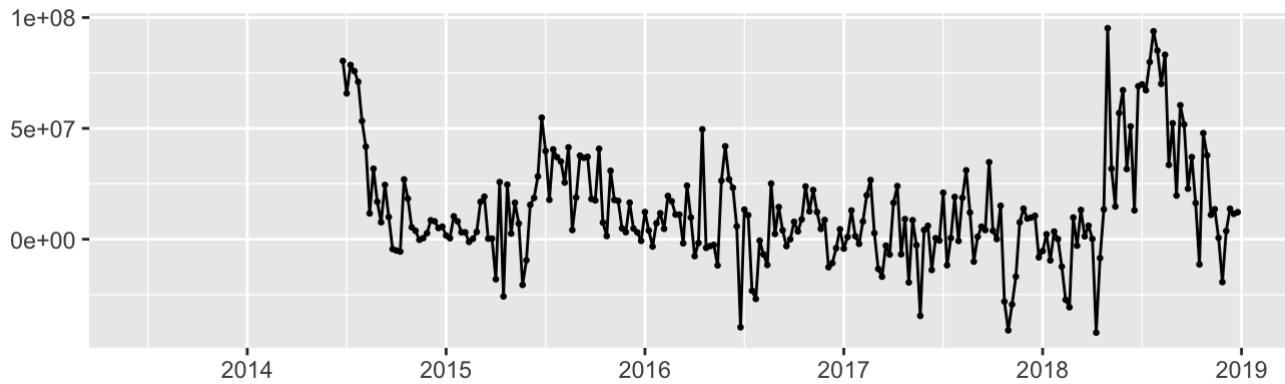
```
model_snaive <- snaive(train_divvy,h=52)

pred_snaive <- model_snaive$mean
```

Check residuals

```
checkresiduals(model_snaive)
```

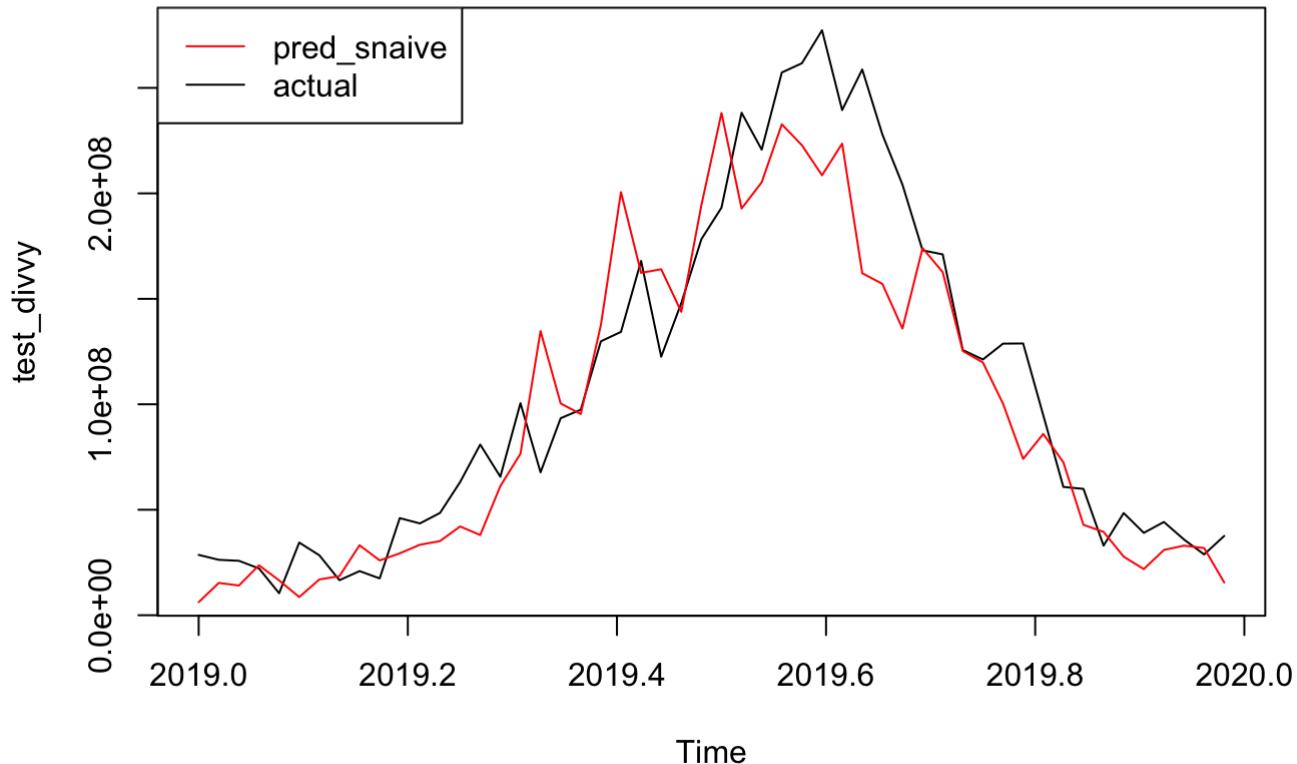
Residuals from Seasonal naive method



```
##
## Ljung-Box test
##
## data: Residuals from Seasonal naive method
## Q* = 452.06, df = 57, p-value < 2.2e-16
##
## Model df: 0. Total lags used: 57
```

Make forecast and compare it with the actual test data

```
plot(test_divvy)
lines(pred_snaive,col='red')
legend('topleft',legend =c('pred_snaive','actual'),col=c('red','black'),lty=1)
```



Calculate MAPE

```
MAPE(pred_snaive,test_divvy)
```

```
## [1] 0.2752298
```

Part 4: Seasonal ARIMA

Build Model with auto.arima

```
arima <- auto.arima(train_divvy,seasonal = 'TRUE', lambda = 'auto',trace=TRUE,approximation = FALSE)
```

```

##                                     : Inf
## ARIMA(2,1,2)(1,1,1)[52]          : 597.6318
## ARIMA(0,1,0)(0,1,0)[52]          : 531.3113
## ARIMA(1,1,0)(1,1,0)[52]          : Inf
## ARIMA(0,1,1)(0,1,1)[52]          : 575.1696
## ARIMA(1,1,0)(0,1,0)[52]          : Inf
## ARIMA(1,1,0)(1,1,1)[52]          : Inf
## ARIMA(1,1,0)(0,1,1)[52]          : 547.5909
## ARIMA(2,1,0)(1,1,0)[52]          : 514.1516
## ARIMA(2,1,0)(0,1,0)[52]          : 559.0914
## ARIMA(2,1,0)(1,1,1)[52]          : Inf
## ARIMA(2,1,0)(0,1,1)[52]          : Inf
## ARIMA(3,1,0)(1,1,0)[52]          : 508.9005
## ARIMA(3,1,0)(0,1,0)[52]          : 557.4164
## ARIMA(3,1,0)(1,1,1)[52]          : Inf
## ARIMA(3,1,0)(0,1,1)[52]          : Inf
## ARIMA(4,1,0)(1,1,0)[52]          : 507.9322
## ARIMA(4,1,0)(0,1,0)[52]          : 552.9106
## ARIMA(4,1,0)(1,1,1)[52]          : Inf
## ARIMA(4,1,0)(0,1,1)[52]          : Inf
## ARIMA(5,1,0)(1,1,0)[52]          : 500.3405
## ARIMA(5,1,0)(0,1,0)[52]          : 540.5245
## ARIMA(5,1,0)(1,1,1)[52]          : Inf
## ARIMA(5,1,0)(0,1,1)[52]          : Inf
## ARIMA(5,1,1)(1,1,0)[52]          : 495.9344
## ARIMA(5,1,1)(0,1,0)[52]          : 532.9529
## ARIMA(5,1,1)(1,1,1)[52]          : Inf
## ARIMA(5,1,1)(0,1,1)[52]          : Inf
## ARIMA(4,1,1)(1,1,0)[52]          : Inf
## ARIMA(5,1,2)(1,1,0)[52]          : 493.8081
## ARIMA(5,1,2)(0,1,0)[52]          : 535.5424
## ARIMA(5,1,2)(1,1,1)[52]          : Inf
## ARIMA(5,1,2)(0,1,1)[52]          : Inf
## ARIMA(4,1,2)(1,1,0)[52]          : Inf
## ARIMA(5,1,3)(1,1,0)[52]          : 492.5909
## ARIMA(5,1,3)(0,1,0)[52]          : Inf
## ARIMA(5,1,3)(1,1,1)[52]          : Inf
## ARIMA(5,1,3)(0,1,1)[52]          : Inf
## ARIMA(4,1,3)(1,1,0)[52]          : 490.1978
## ARIMA(4,1,3)(0,1,0)[52]          : 536.9072
## ARIMA(4,1,3)(1,1,1)[52]          : Inf
## ARIMA(4,1,3)(0,1,1)[52]          : Inf
## ARIMA(3,1,3)(1,1,0)[52]          : Inf
## ARIMA(4,1,4)(1,1,0)[52]          : 493.2381
## ARIMA(3,1,2)(1,1,0)[52]          : Inf
## ARIMA(3,1,4)(1,1,0)[52]          : Inf
## ARIMA(5,1,4)(1,1,0)[52]          : Inf
##
## Best model: ARIMA(4,1,3)(1,1,0)[52]

```

arima

```

## Series: train_divvy
## ARIMA(4,1,3)(1,1,0)[52]
## Box Cox transformation: lambda= 0.03488702
##
## Coefficients:
##             ar1      ar2      ar3      ar4      ma1      ma2      ma3      sar1
##             1.1765   -1.2798   0.3741  -0.1505  -1.6976   1.6812  -0.7769  -0.4566
## s.e.     0.1293    0.1099   0.1202    0.0965   0.1315   0.1210   0.1363   0.0733
##
## sigma^2 estimated as 0.427: log likelihood=-235.7
## AIC=489.39   AICc=490.2   BIC=520.49

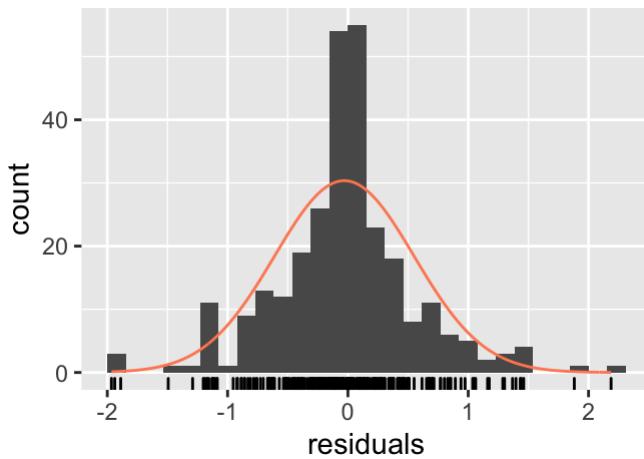
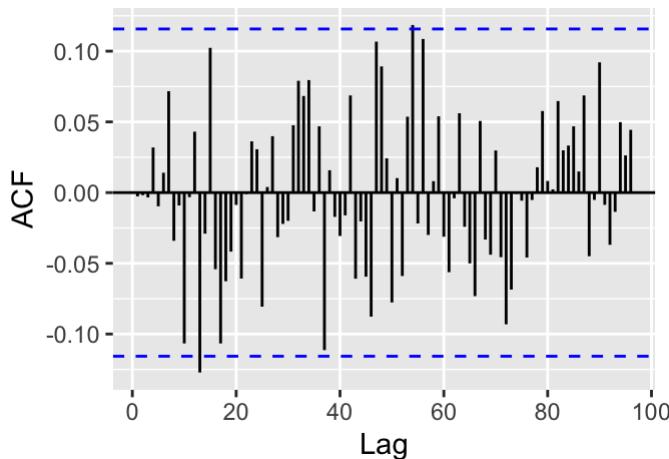
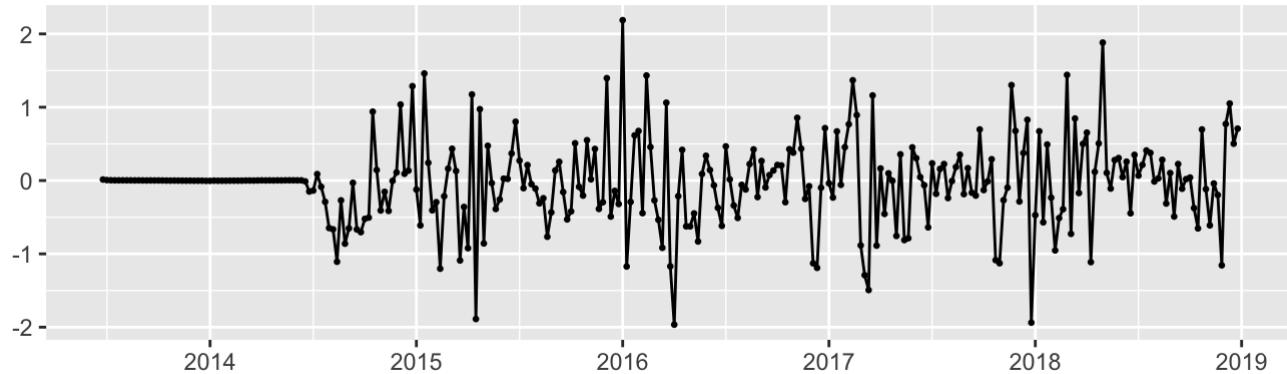
```

(4,1,3) (1,1,0) period=52

Check residuals

```
checkresiduals(arima)
```

Residuals from ARIMA(4,1,3)(1,1,0)[52]



```

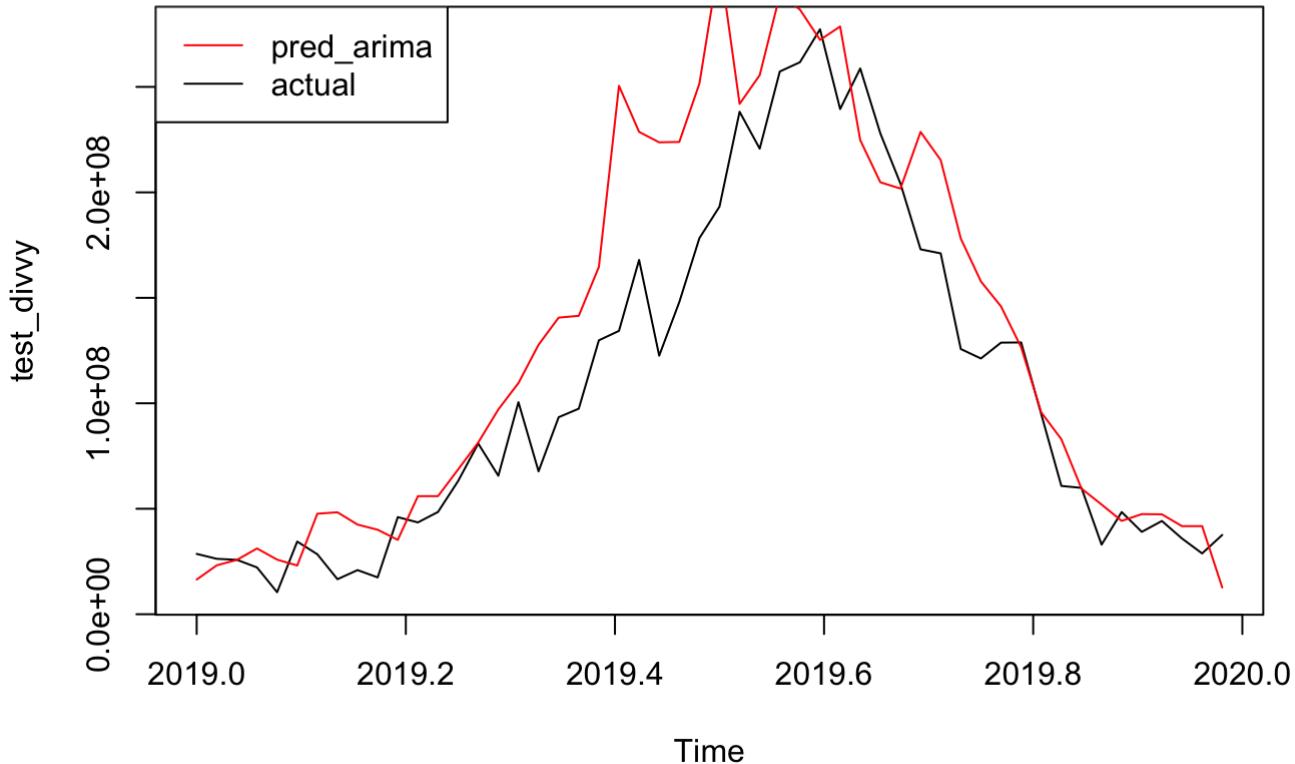
##
## Ljung-Box test
##
## data: Residuals from ARIMA(4,1,3)(1,1,0)[52]
## Q* = 65.523, df = 49, p-value = 0.05738
##
## Model df: 8. Total lags used: 57

```

P=0.0574; Barely not reject the null hypothesis

Make Predictions and Compare with the actual test data

```
pred_arima <- forecast(arima,h=52)
plot(test_divvy)
lines(pred_arima$mean,col='red')
legend('topleft',legend =c('pred_arima','actual'),col=c('red','black'),lty=1)
```



Calculate MAPE

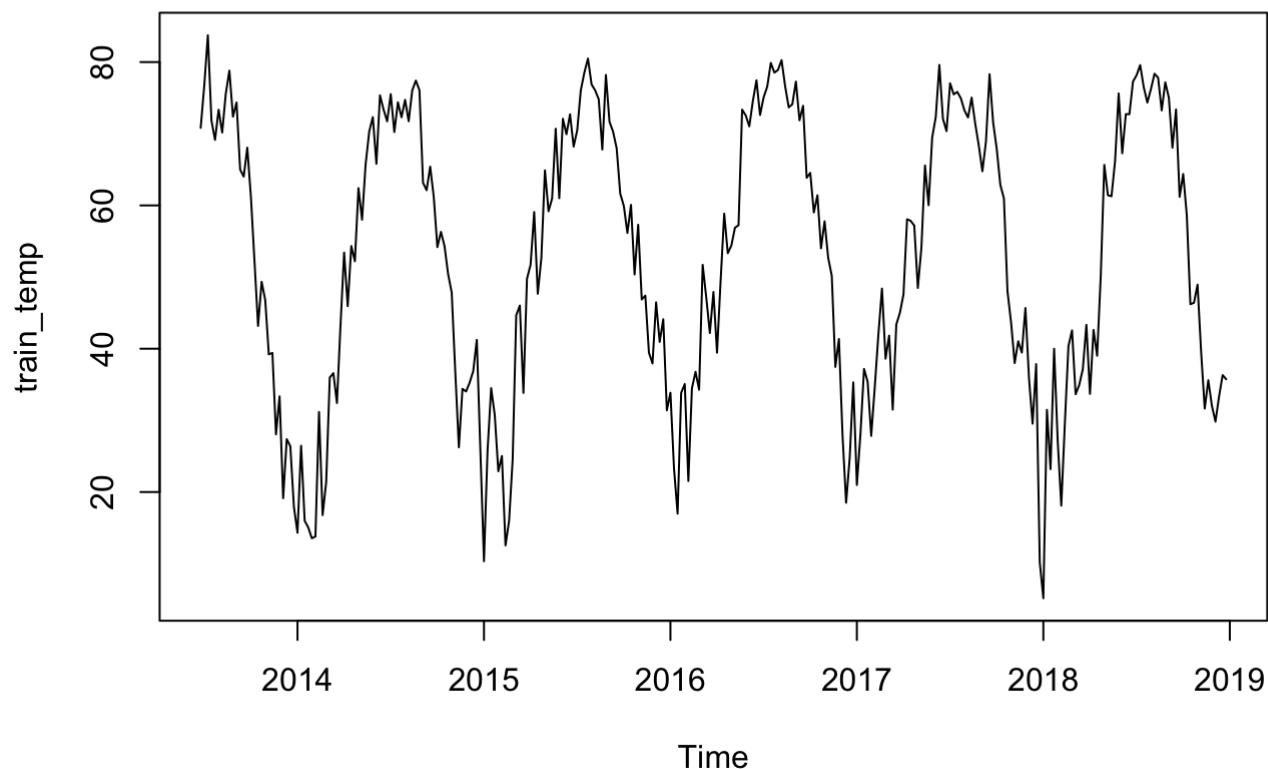
```
MAPE(pred_arima$mean,test_divvy)
```

```
## [1] 0.376772
```

Part 5: Regression with ARIMA error (with Temp)

Plot temp data and see if Boxcox is needed

```
plot(train_temp)
```

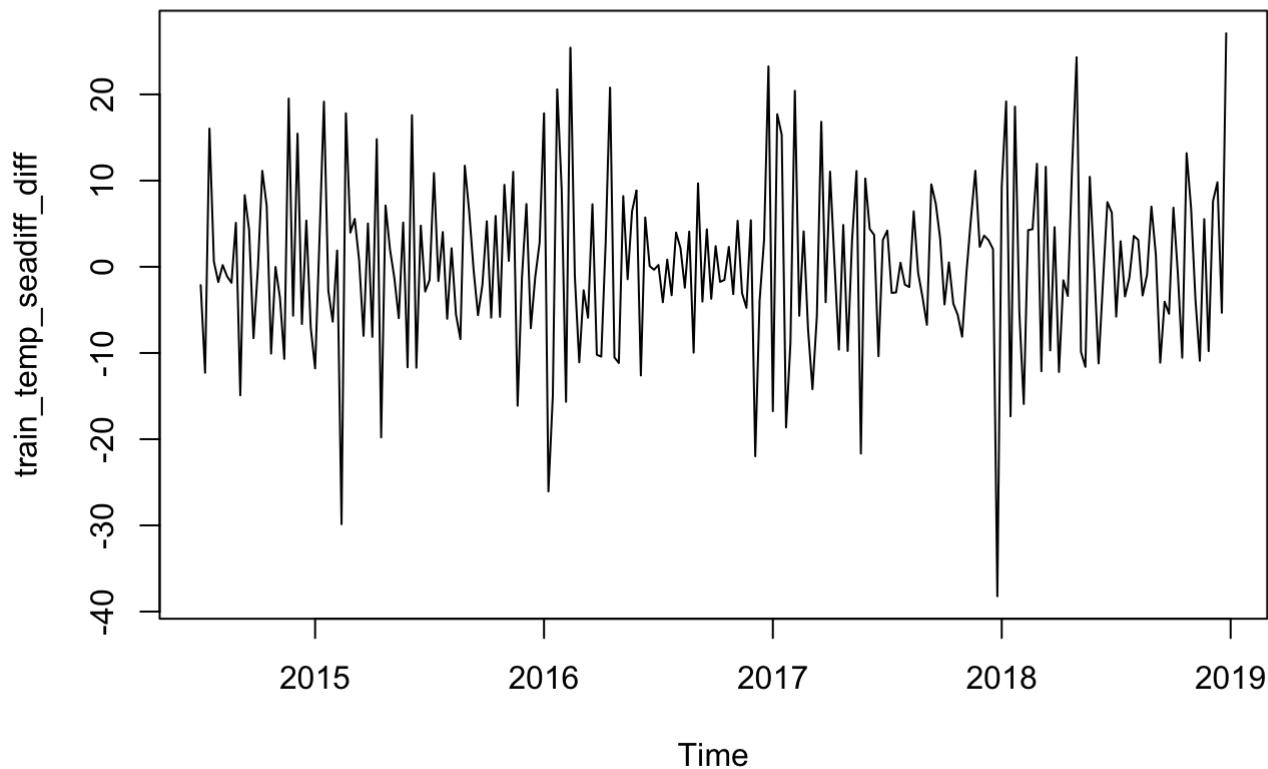


No need for BoxCox

Use Seasonal Diff and First Order Diff to remove and Seasonality and make it trend stationary

```
train_temp_seadiff_diff <- diff(diff(train_temp,lag=52))

plot(train_temp_seadiff_diff)
```



```
kpss.test(train_temp_seadiff_diff)
```

```
## Warning in kpss.test(train_temp_seadiff_diff): p-value greater than printed p-
## value
```

```
##
## KPSS Test for Level Stationarity
##
## data: train_temp_seadiff_diff
## KPSS Level = 0.039703, Truncation lag parameter = 4, p-value = 0.1
```

larger p indicating trend is stationary

Prepare new training data for model fitting

```
train_divvy_new <- train_divvy[54:287]

train_divvy_new <- ts(train_divvy_new,frequency = 52,start=c(2014,27))
```

Fit model using xreg = train_temp_seadiff_diff (stationary)

```
model_arimae <- auto.arima(train_divvy_new,xreg=train_temp_seadiff_diff,lambda = 0.0349,
d=1,D=1)
model_arimae
```

```

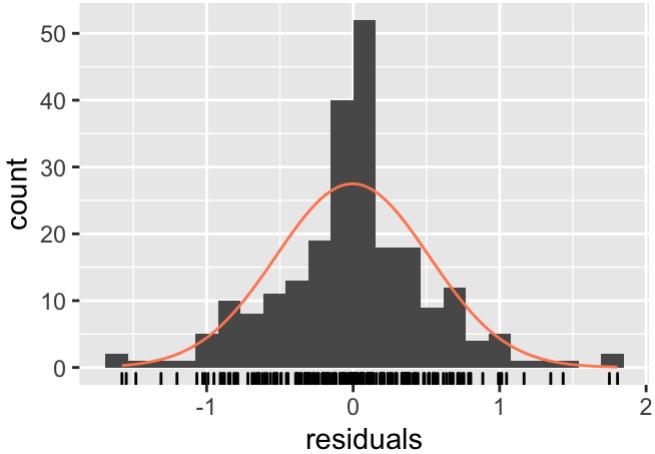
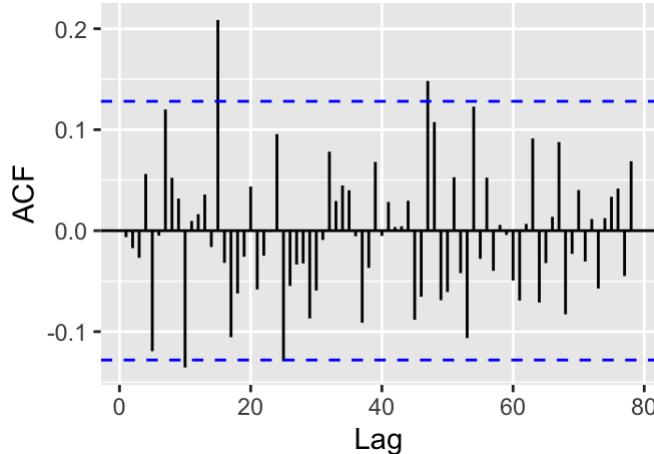
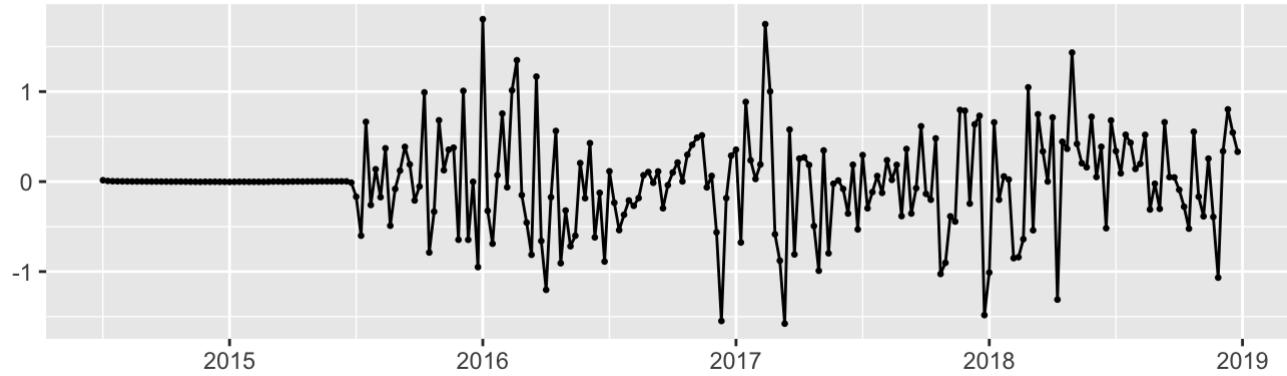
## Series: train_divvy_new
## Regression with ARIMA(0,1,3)(1,1,0)[52] errors
## Box Cox transformation: lambda= 0.0349
##
## Coefficients:
##             ma1      ma2      ma3     sar1     xreg
##             -0.4921  -0.3048  -0.0695  -0.3671  0.0161
## s.e.      0.0775   0.0738   0.0732   0.0853  0.0025
##
## sigma^2 estimated as 0.3611: log likelihood=-166.44
## AIC=344.87    AICc=345.35    BIC=364.06

```

Check Residuals

```
checkresiduals(model_arimae)
```

Residuals from Regression with ARIMA(0,1,3)(1,1,0)[52] errors



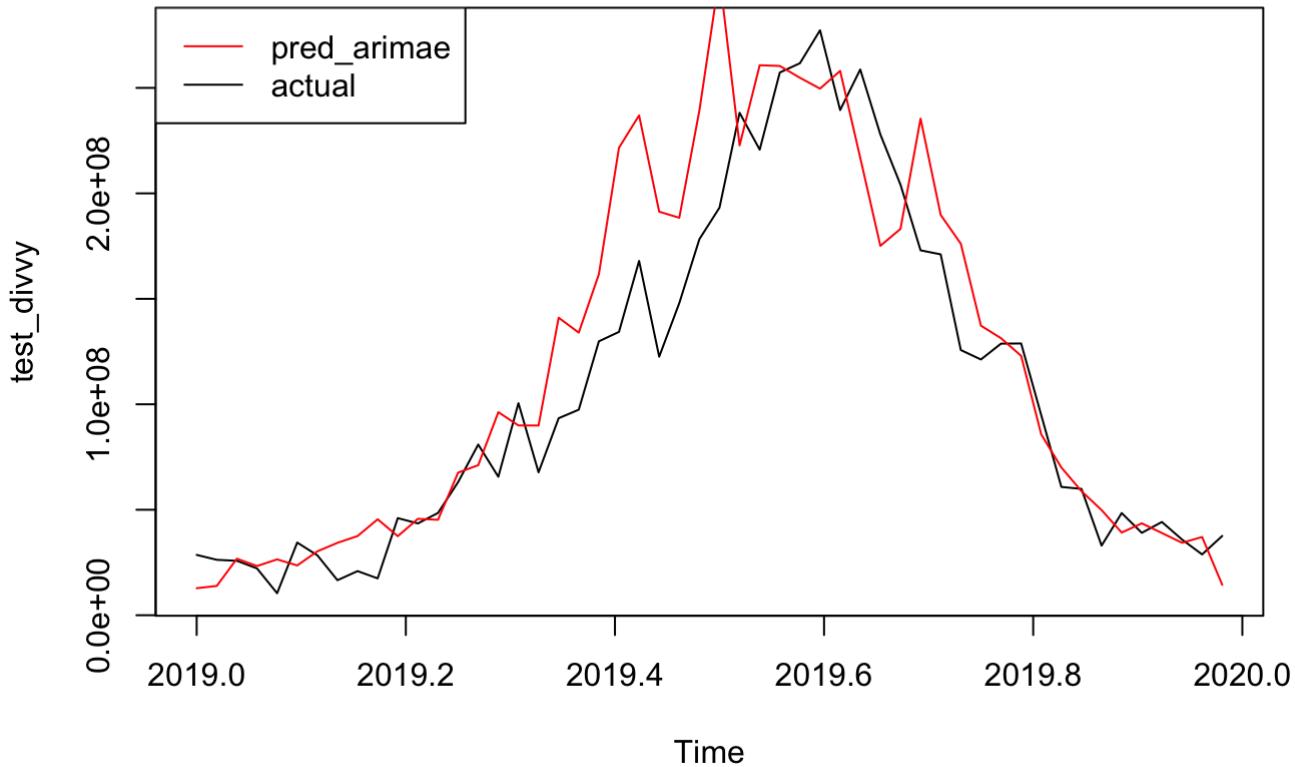
```

##
## Ljung-Box test
##
## data: Residuals from Regression with ARIMA(0,1,3)(1,1,0)[52] errors
## Q* = 59.302, df = 42, p-value = 0.04025
##
## Model df: 5. Total lags used: 47

```

Make Predictions and Compare with the actual test data

```
test_temp_new <- c(train_temp[235:287],test_temp)
pred_arimae <- forecast(model_arimae,xreg=diff(diff(test_temp_new,lag=52)),h=52)
plot(test_divvy)
lines(pred_arimae$mean,col='red')
legend('topleft',legend =c('pred_arimae','actual'),col=c('red','black'),lty=1)
```



Calculate MAPE

```
MAPE(pred_arimae$mean,test_divvy)
```

```
## [1] 0.3077718
```

Part 6: Vector AutoRegression (VAR)

```
library(vars)
data = cbind(divvy=train_divvy,temp=train_temp)
```

Check what lag order is appropriate for VAR model

```
VARselect(data,lag.max=5,type='both')$selection
```

```
## AIC(n)  HQ(n)  SC(n) FPE(n)
##      3       3       1       3
```

Fit model with lag = 1 and check residual independence

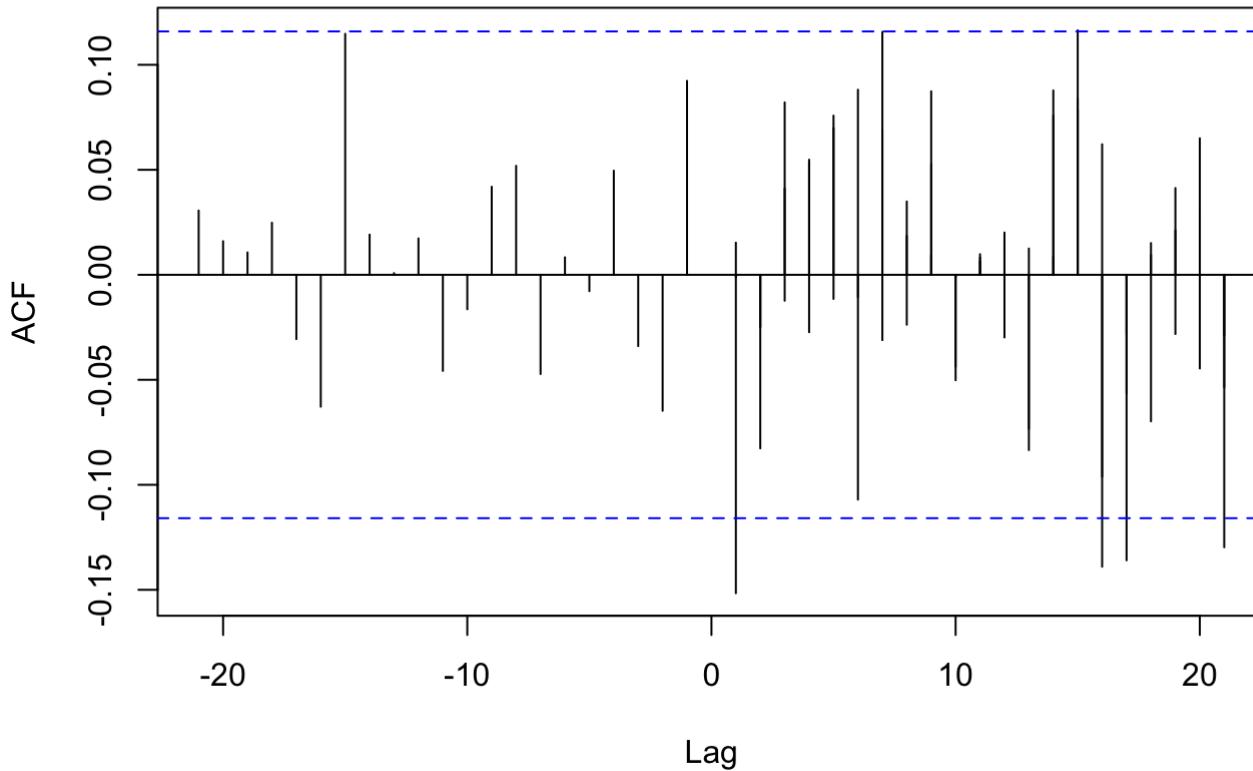
```
var1 <- VAR(data,p=1,type='both',season = 52)

serial.test(var1,lags.pt=10,type='PT.asymptotic')
```

```
##
## Portmanteau Test (asymptotic)
##
## data: Residuals of VAR object var1
## Chi-squared = 52.23, df = 36, p-value = 0.03927
```

```
acf(residuals(var1))
```

Series residuals(var1)



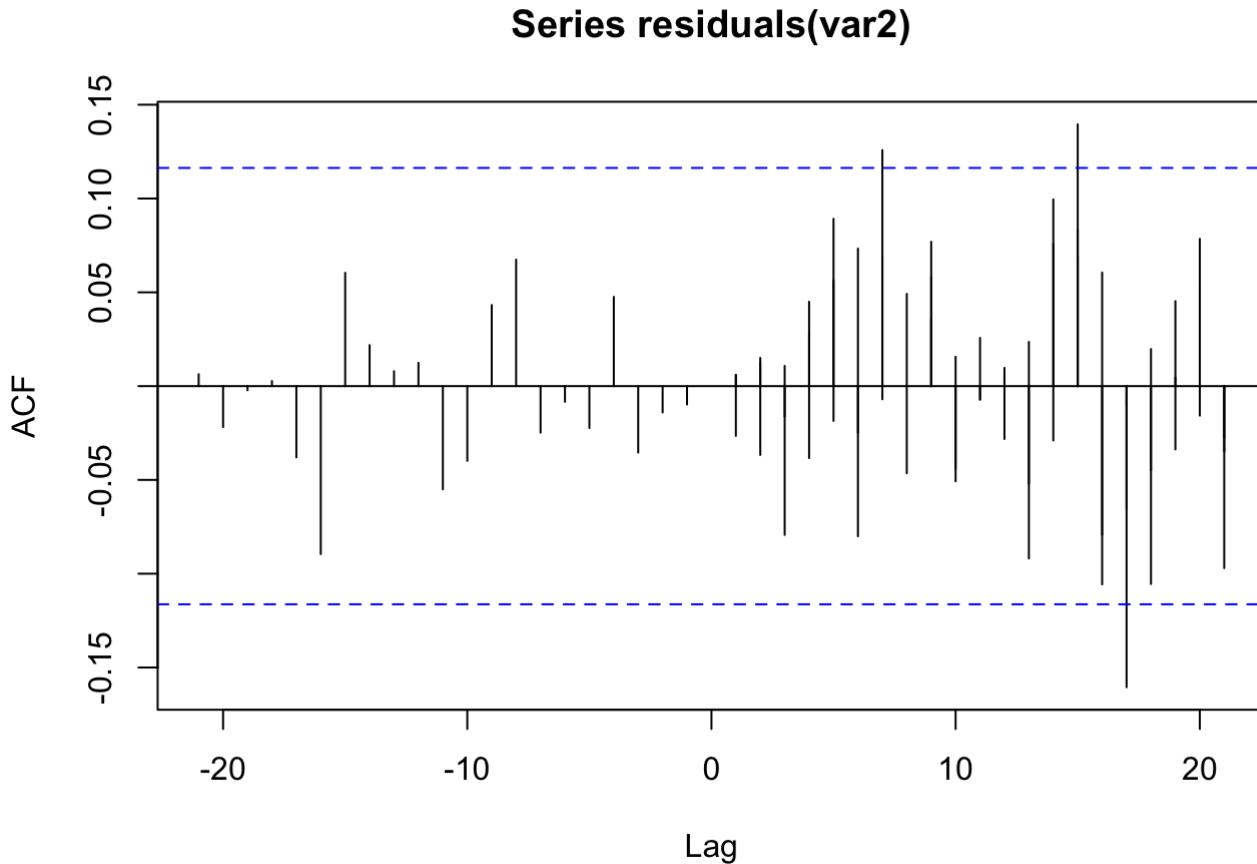
Fit model with lag = 3 and check residual independence

```
var2 <- VAR(data,p=3,type='both',season = 52)

serial.test(var2,lags.pt=10,type='PT.asymptotic')
```

```
##  
## Portmanteau Test (asymptotic)  
##  
## data: Residuals of VAR object var2  
## Chi-squared = 30.372, df = 28, p-value = 0.3456
```

```
acf(residuals(var2))
```



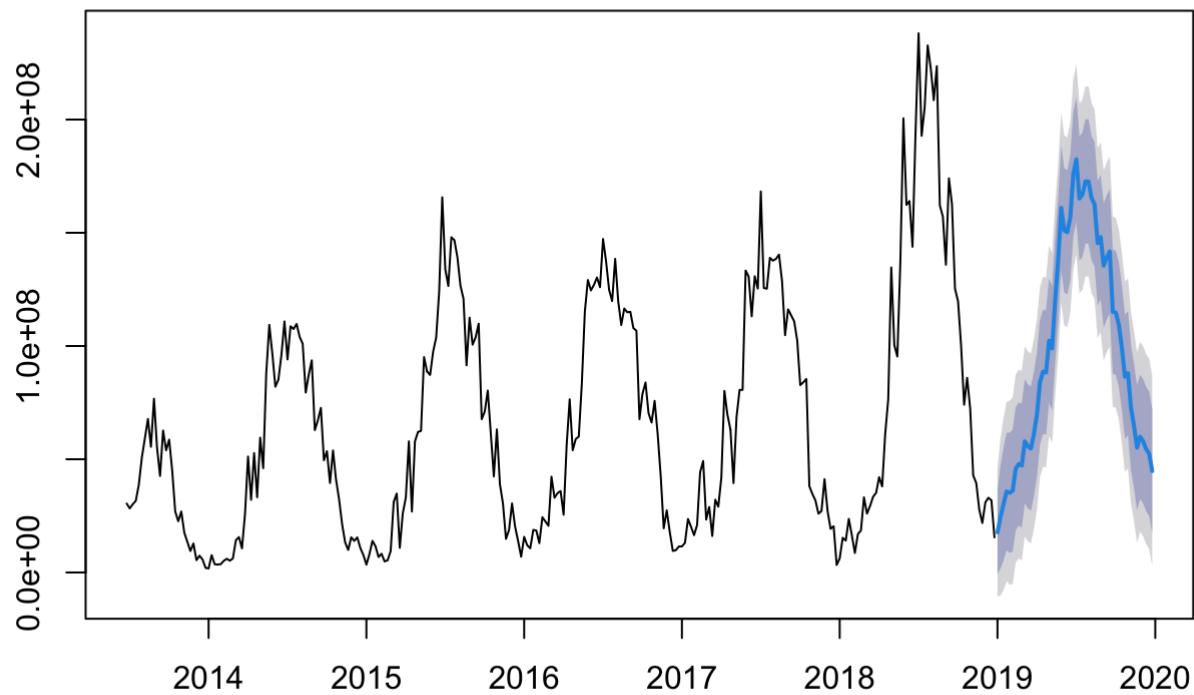
In both cases, residuals are not white noise due to small p values from serial test and spike in acf.

Make predictions and Compare with the actual test data for both models

```
pred_var1 <- forecast(var1,h=52)  
pred_var2 <- forecast(var2,h=52)
```

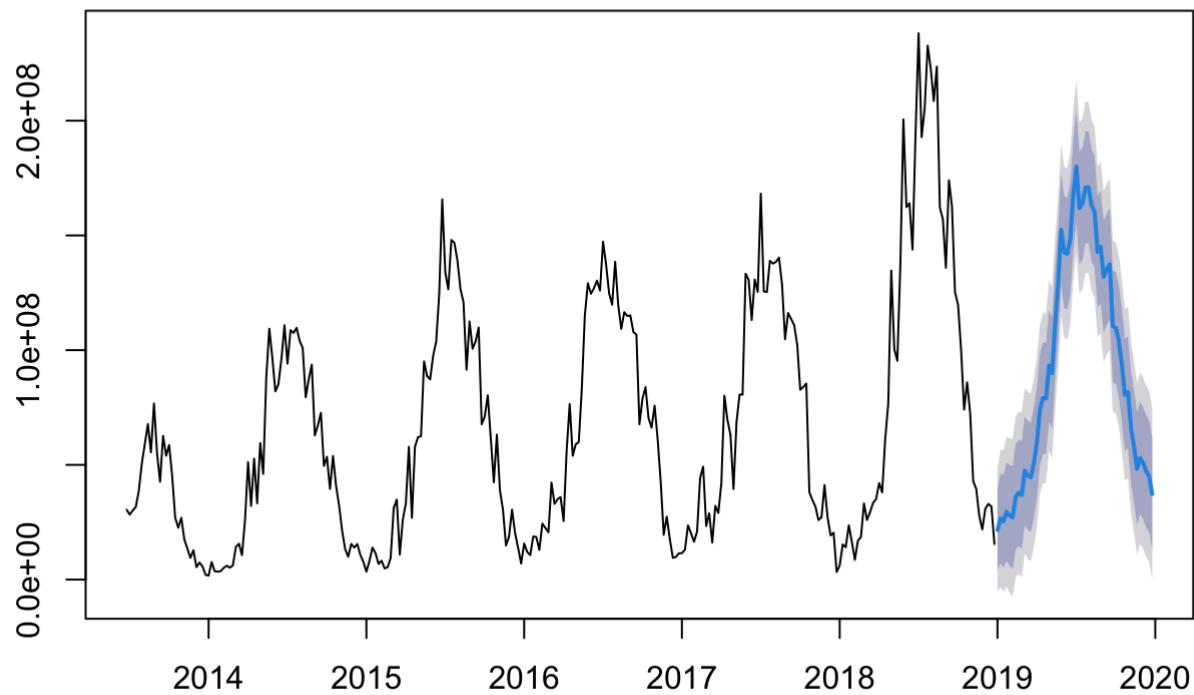
```
plot(pred_var1$forecast$divvyy)
```

Forecasts from VAR(1)



```
plot(pred_var2$forecast$divvy)
```

Forecasts from VAR(3)



```
MAPE(pred_var1$forecast$divvy$mean,test_divvy)
```

```
## [1] 0.3989797
```

```
MAPE(pred_var2$forecast$divvy$mean,test_divvy)
```

```
## [1] 0.2948619
```

Part 7: Fourier Transform

Let's first check the periodogram of train_divvy dataset.

```
kpss.test(train_divvy_boxcox)
```

```
##
## KPSS Test for Level Stationarity
##
## data: train_divvy_boxcox
## KPSS Level = 0.59309, Truncation lag parameter = 5, p-value = 0.02326
```

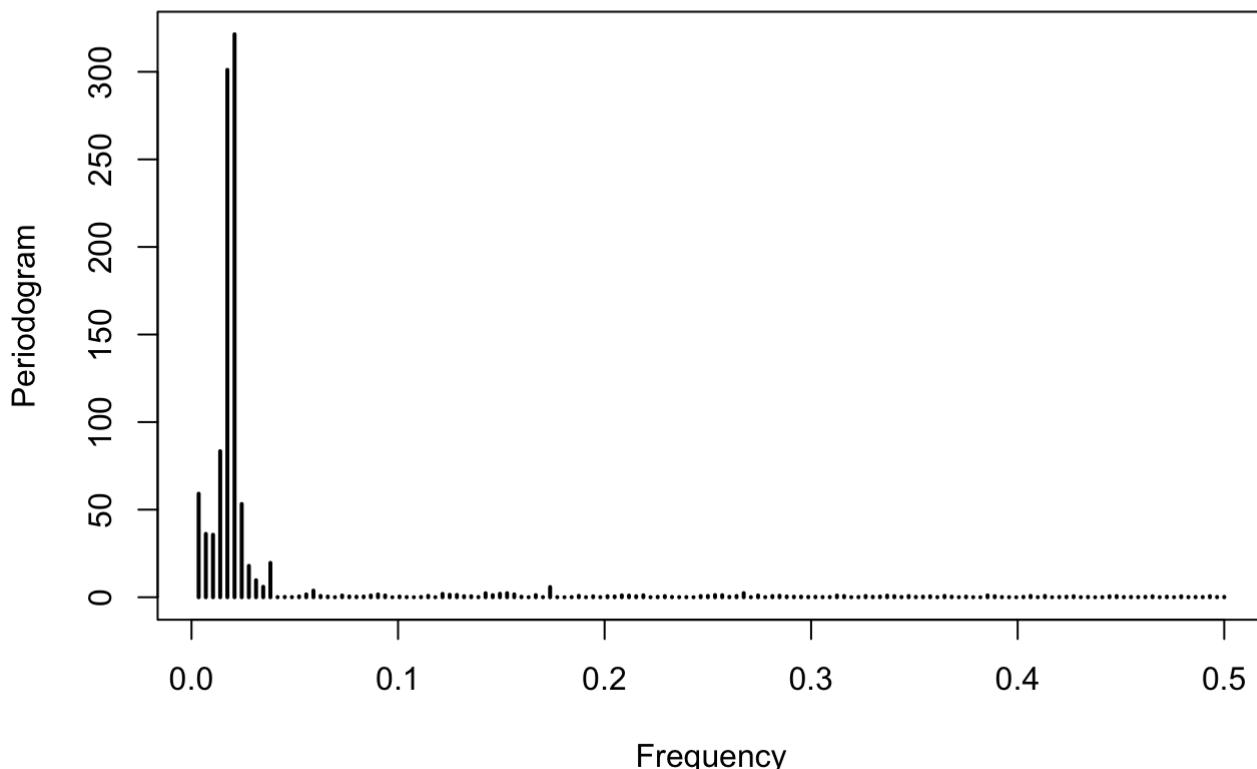
```
adf.test(train_divvy_boxcox)
```

```
##  
## Augmented Dickey-Fuller Test  
##  
## data: train_divvy_boxcox  
## Dickey-Fuller = -3.7341, Lag order = 6, p-value = 0.02273  
## alternative hypothesis: stationary
```

The Kpss test shows that the series is trend stationary ($p>0.05$).

The adf test shows that the data is stationary ($p<0.05$).

```
periodogram(train_divvy_boxcox)  
temp <- periodogram(train_divvy_boxcox)
```



```
temp$freq
```

```
## [1] 0.003472222 0.006944444 0.010416667 0.013888889 0.017361111 0.020833333
## [7] 0.024305556 0.027777778 0.031250000 0.034722222 0.038194444 0.041666667
## [13] 0.045138889 0.048611111 0.052083333 0.055555556 0.059027778 0.062500000
## [19] 0.065972222 0.069444444 0.072916667 0.076388889 0.079861111 0.083333333
## [25] 0.086805556 0.090277778 0.093750000 0.097222222 0.100694444 0.104166667
## [31] 0.107638889 0.111111111 0.114583333 0.118055556 0.121527778 0.125000000
## [37] 0.128472222 0.131944444 0.135416667 0.138888889 0.142361111 0.145833333
## [43] 0.149305556 0.152777778 0.156250000 0.159722222 0.163194444 0.166666667
## [49] 0.170138889 0.173611111 0.177083333 0.180555556 0.184027778 0.187500000
## [55] 0.190972222 0.194444444 0.197916667 0.201388889 0.204861111 0.208333333
## [61] 0.211805556 0.215277778 0.218750000 0.222222222 0.225694444 0.229166667
## [67] 0.232638889 0.236111111 0.239583333 0.243055556 0.246527778 0.250000000
## [73] 0.253472222 0.256944444 0.260416667 0.263888889 0.267361111 0.270833333
## [79] 0.274305556 0.277777778 0.281250000 0.284722222 0.288194444 0.291666667
## [85] 0.295138889 0.298611111 0.302083333 0.305555556 0.309027778 0.312500000
## [91] 0.315972222 0.319444444 0.322916667 0.326388889 0.329861111 0.333333333
## [97] 0.336805556 0.340277778 0.343750000 0.347222222 0.350694444 0.354166667
## [103] 0.357638889 0.361111111 0.364583333 0.368055556 0.371527778 0.375000000
## [109] 0.378472222 0.381944444 0.385416667 0.388888889 0.392361111 0.395833333
## [115] 0.399305556 0.402777778 0.406250000 0.409722222 0.413194444 0.416666667
## [121] 0.420138889 0.423611111 0.427083333 0.430555556 0.434027778 0.437500000
## [127] 0.440972222 0.444444444 0.447916667 0.451388889 0.454861111 0.458333333
## [133] 0.461805556 0.465277778 0.468750000 0.472222222 0.475694444 0.479166667
## [139] 0.482638889 0.486111111 0.489583333 0.493055556 0.496527778 0.500000000
```

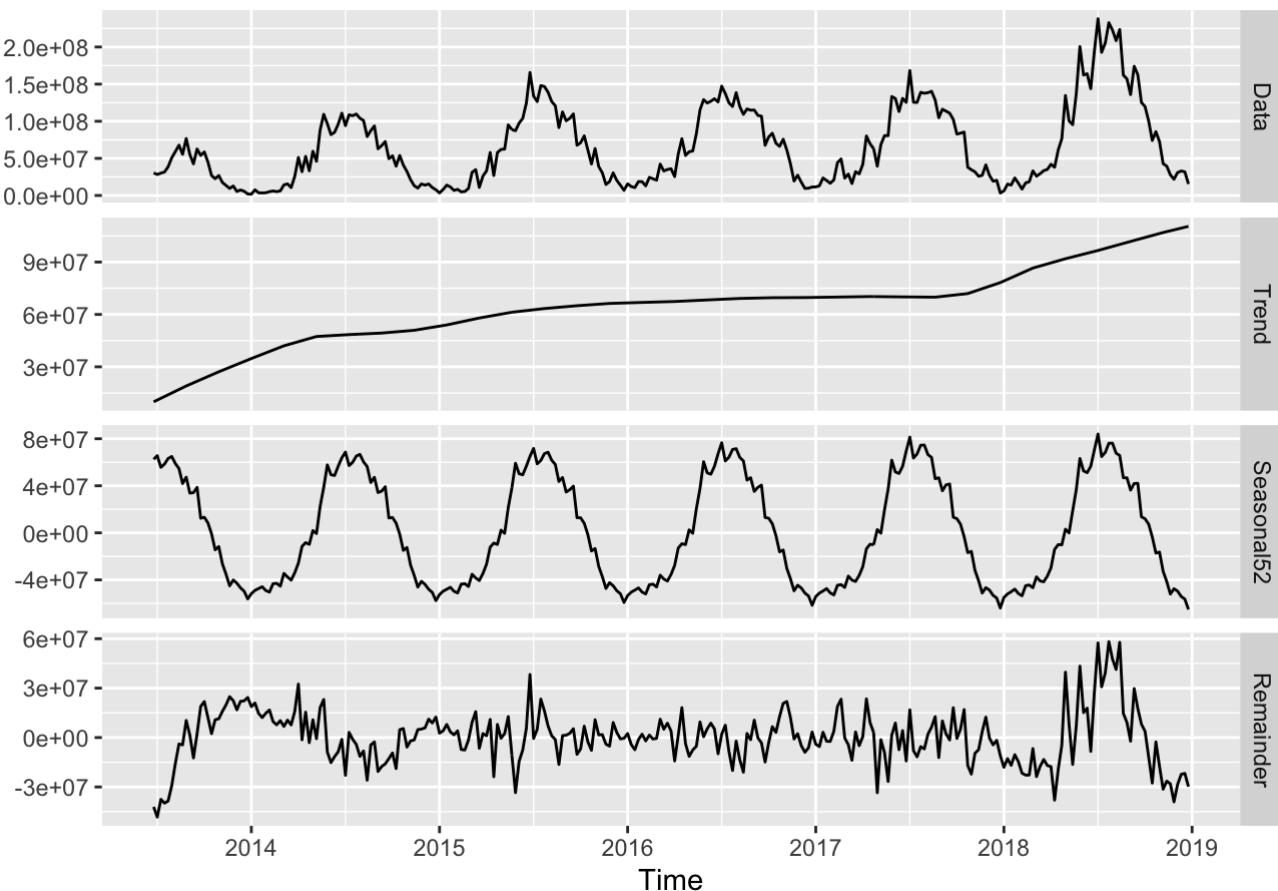
```
Max_freq <- temp$freq[which.max(temp$spec) ]
1/Max_freq
```

```
## [1] 48
```

The highest two periodogram fall on frequency 0.017361111 and 0.020833333, which is 57.6 weeks and 48 weeks.

Plot STL to show the important seasonality in this data.

```
autoplot(mstl(train_divvy))
```



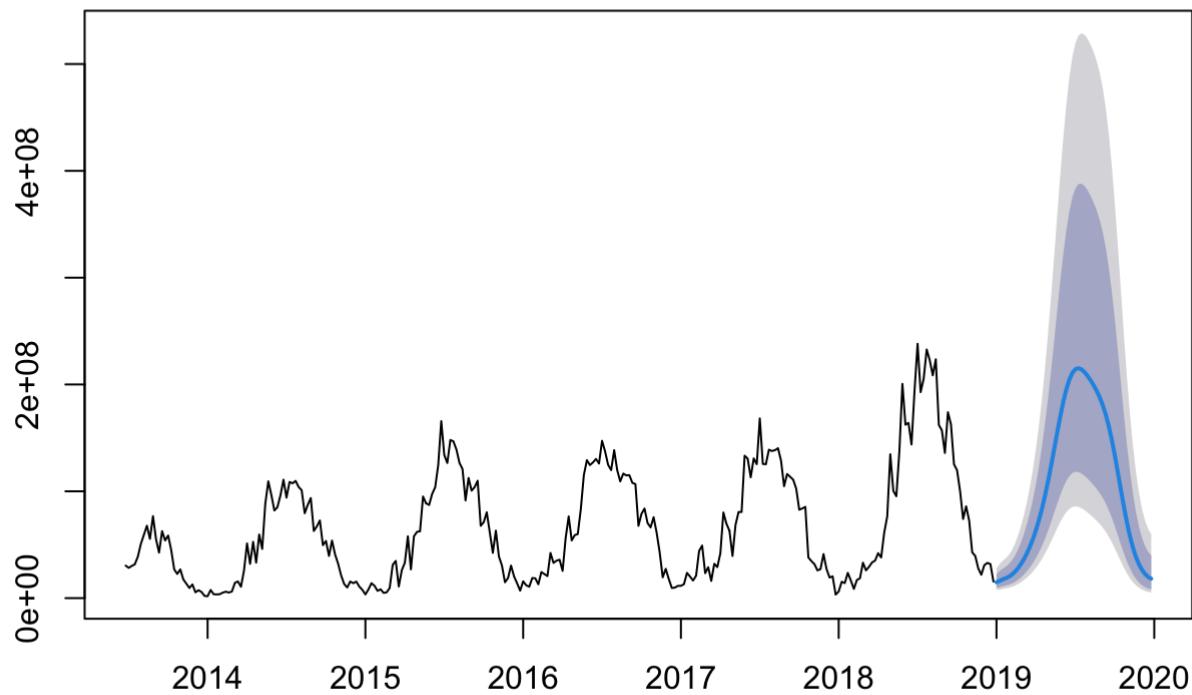
The 52 period is the most significant seasonality in this dataset.

Combine Fourier terms with ARIMA errors

K=3

```
arima_fourier_3 <- auto.arima(train_divvy, xreg=fourier(train_divvy,3), seasonal=FALSE, 1
lambda = "auto")
arima_fourier_forecast_3 <- forecast(arima_fourier_3, xreg=fourier(train_divvy, 3, 52))
plot(arima_fourier_forecast_3)
```

Forecasts from Regression with ARIMA(2,1,2) errors

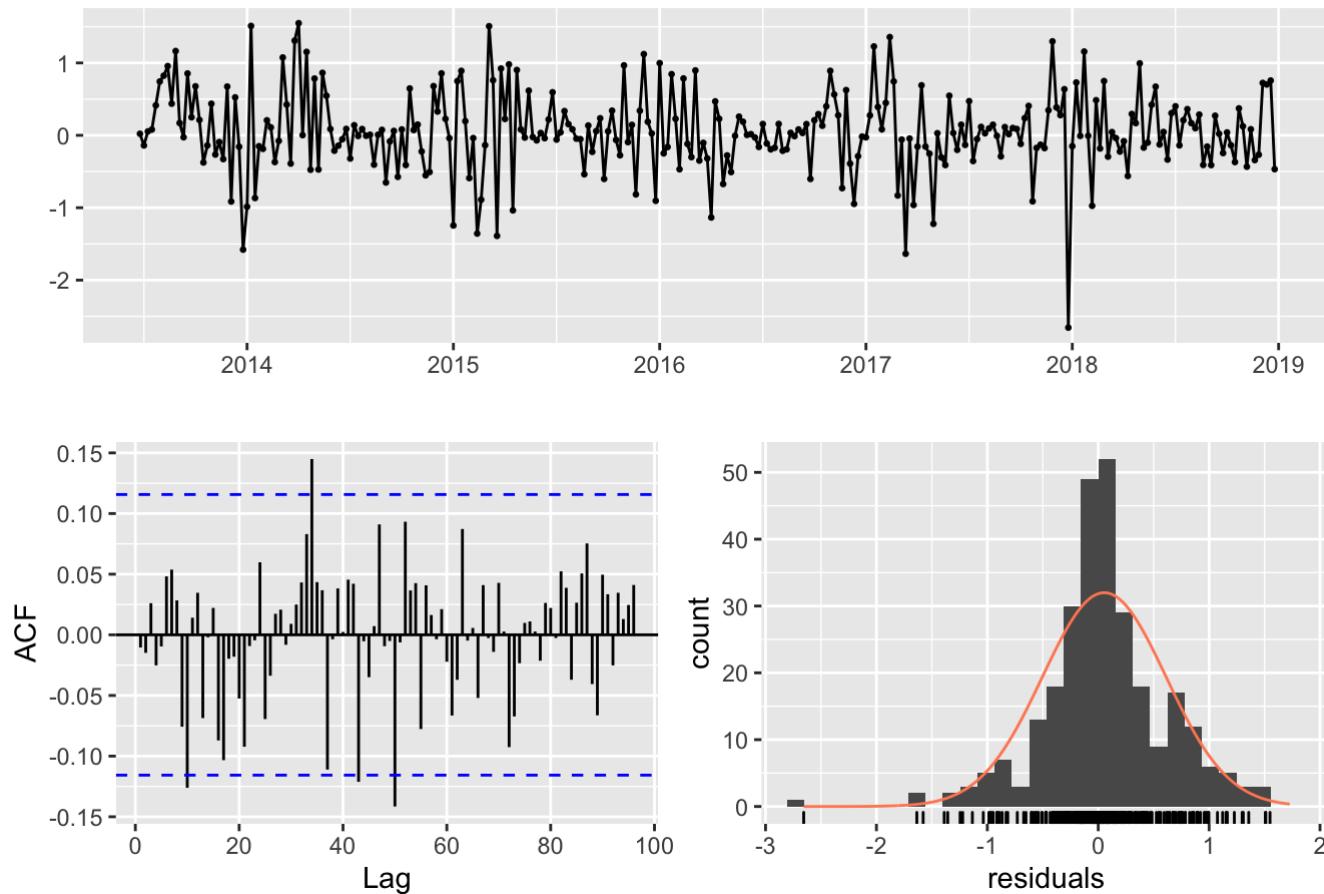


```
summary(arima_fourier_3)
```

```
## Series: train_divvy
## Regression with ARIMA(2,1,2) errors
## Box Cox transformation: lambda= 0.03488702
##
## Coefficients:
##             ar1      ar2      ma1      ma2    S1-52    C1-52    S2-52    C2-52
##             0.6597 -0.2782 -1.1942  0.3367  1.1340  2.0534 -0.2436 -0.2874
## s.e.     0.3661  0.1655  0.3649  0.3575  0.1018  0.0979  0.0635  0.0610
##             S3-52    C3-52
##             -0.0113  0.1291
## s.e.     0.0539  0.0532
##
## sigma^2 estimated as 0.325: log likelihood=-240.48
## AIC=502.96   AICc=503.92   BIC=543.18
##
## Training set error measures:
##               ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
## Training set 2208815 14009139 9723598 -2.305292 22.59074 0.5079432 0.003965753
```

```
checkresiduals(arima_fourier_3)
```

Residuals from Regression with ARIMA(2,1,2) errors

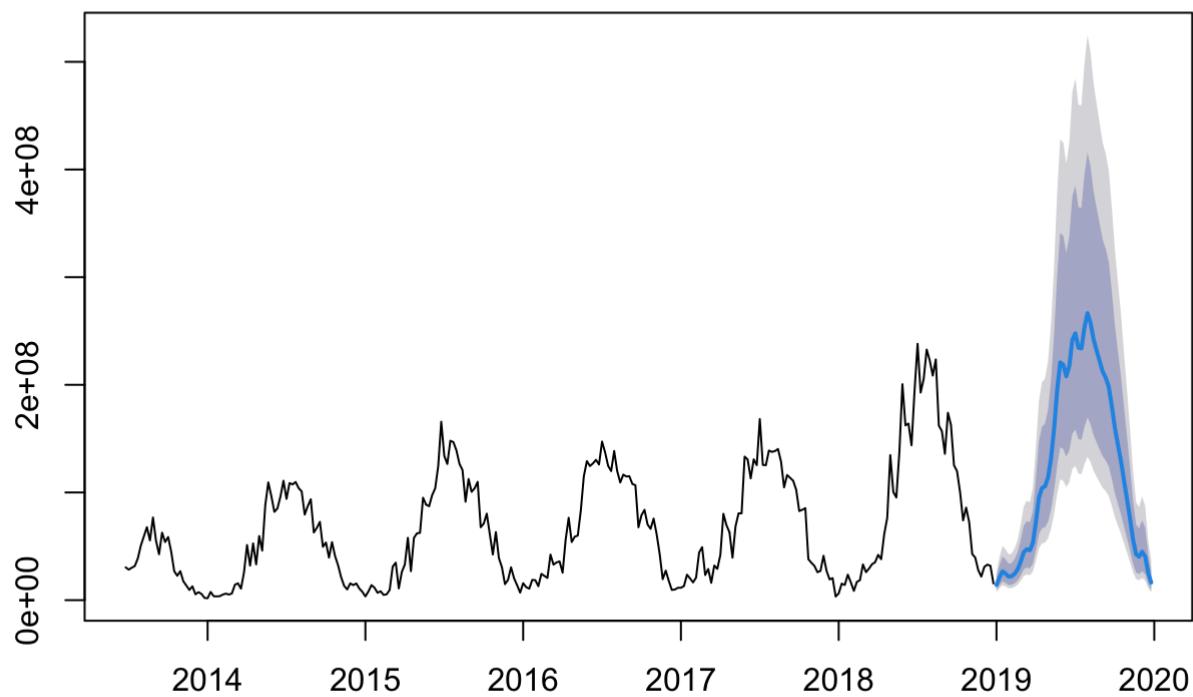


```
##  
## Ljung-Box test  
##  
## data: Residuals from Regression with ARIMA(2,1,2) errors  
## Q* = 62.722, df = 47, p-value = 0.06219  
##  
## Model df: 10. Total lags used: 57
```

K=13

```
arima_fourier_13 <- auto.arima(train_divvy, xreg=fourier(train_divvy,13), seasonal=FALSE  
,lambda = "auto")  
arima_fourier_forecast_13 <- forecast(arima_fourier_13, xreg=fourier(train_divvy, 13, 52  
))  
plot(arima_fourier_forecast_13)
```

Forecasts from Regression with ARIMA(3,1,1) errors



The forecast now looks more reasonable.

```
summary(arima_fourier_13)
```

```

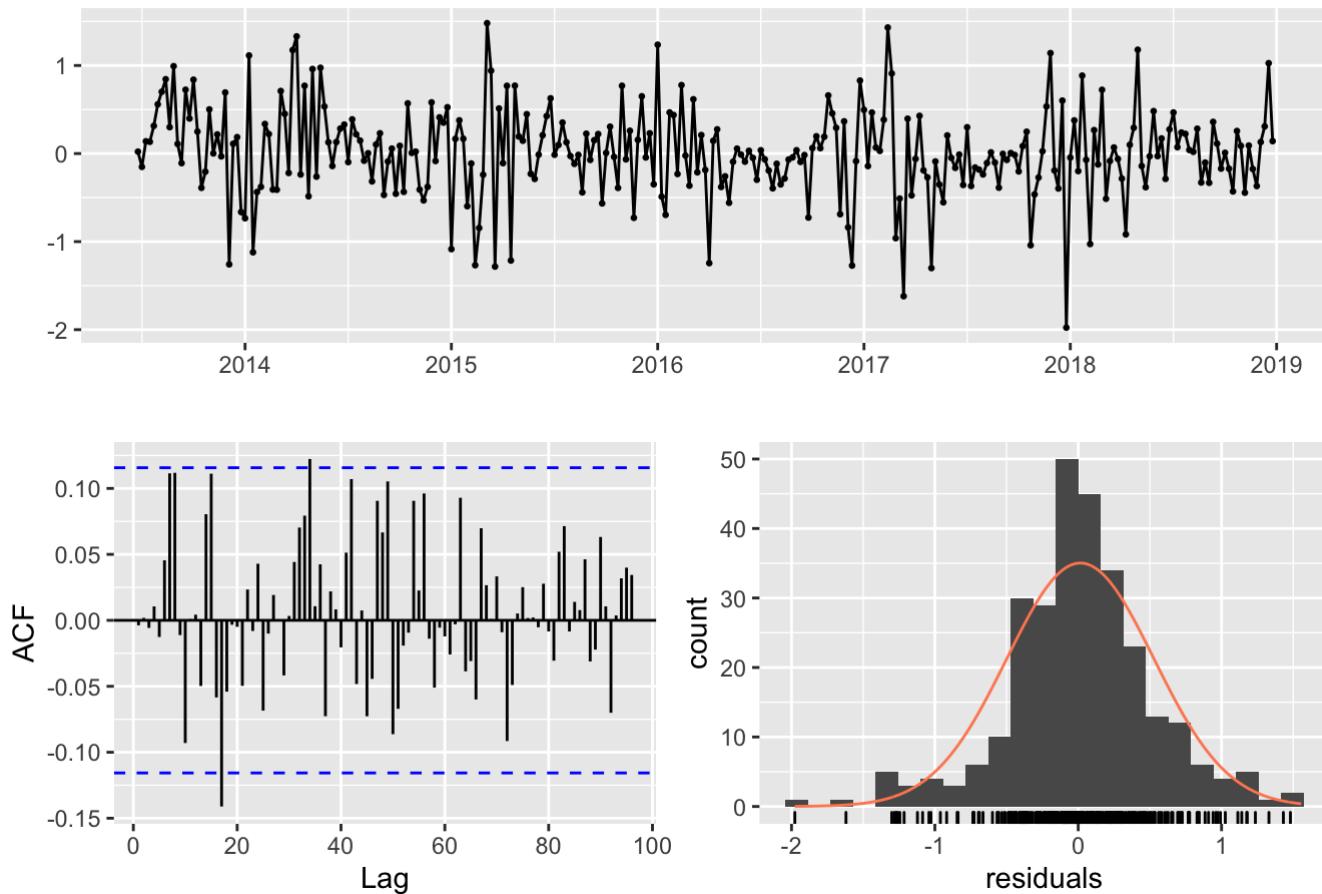
## Series: train_divvy
## Regression with ARIMA(3,1,1) errors
## Box Cox transformation: lambda= 0.03488702
##
## Coefficients:
##             ar1      ar2      ar3      ma1     drift    S1-52     C1-52     S2-52     C2-52
##             0.3285  0.0057  0.1483 -0.9470  0.0097  1.1023  2.0470 -0.2467 -0.2912
## s.e.   0.0668  0.0654  0.0646  0.0306  0.0034  0.0874  0.0852  0.0732  0.0721
##             S3-52     C3-52     S4-52     C4-52     S5-52     C5-52     S6-52     C6-52
##            -0.0261  0.1224 -0.0162 -0.049  -0.0047  0.0166 -0.0900 -0.0254
## s.e.   0.0636  0.0632  0.0563  0.056  0.0506  0.0505  0.0465  0.0465
##             S7-52     C7-52     S8-52     C8-52     S9-52     C9-52     S10-52    C10-52
##            0.0490  0.0239 -0.1431 -0.0648  0.2051 -0.0368 -0.0480  0.0248
## s.e.   0.0436  0.0436  0.0417  0.0417  0.0406  0.0406  0.0401  0.0402
##             S11-52    C11-52    S12-52    C12-52    S13-52    C13-52
##            0.0866  0.0074 -0.0595 -0.0135  0.0617 -0.0348
## s.e.   0.0401  0.0402  0.0405  0.0406  0.0409  0.0410
##
## sigma^2 estimated as 0.2963: log likelihood=-216.1
## AIC=496.21  AICc=504.56  BIC=613.2
##
## Training set error measures:
##               ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
## Training set 945492 13595850 9472912 -3.537875 21.39672 0.4948478 0.1057392

```

Let's check the residual of the model

```
checkresiduals(arima_fourier_13)
```

Residuals from Regression with ARIMA(3,1,1) errors

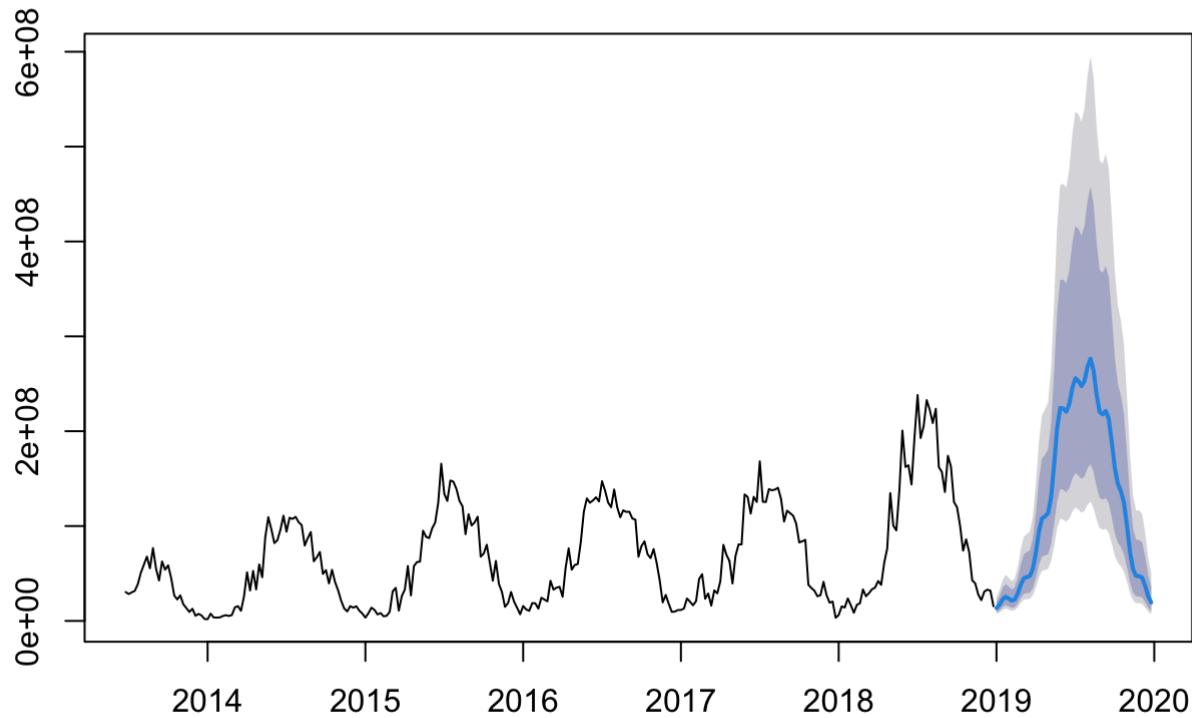


```
##  
## Ljung-Box test  
##  
## data: Residuals from Regression with ARIMA(3,1,1) errors  
## Q* = 68.077, df = 26, p-value = 1.259e-05  
##  
## Model df: 31. Total lags used: 57
```

Let's tune the parameter K=9.

```
arima_fourier_9 <- auto.arima(train_divvy, xreg=fourier(train_divvy, 9), seasonal=FALSE, 1  
lambda = "auto")  
arima_fourier_forecast_9 <- forecast(arima_fourier_9, xreg=fourier(train_divvy, 9, 52))  
plot(arima_fourier_forecast_9)
```

Forecasts from Regression with ARIMA(0,1,2) errors

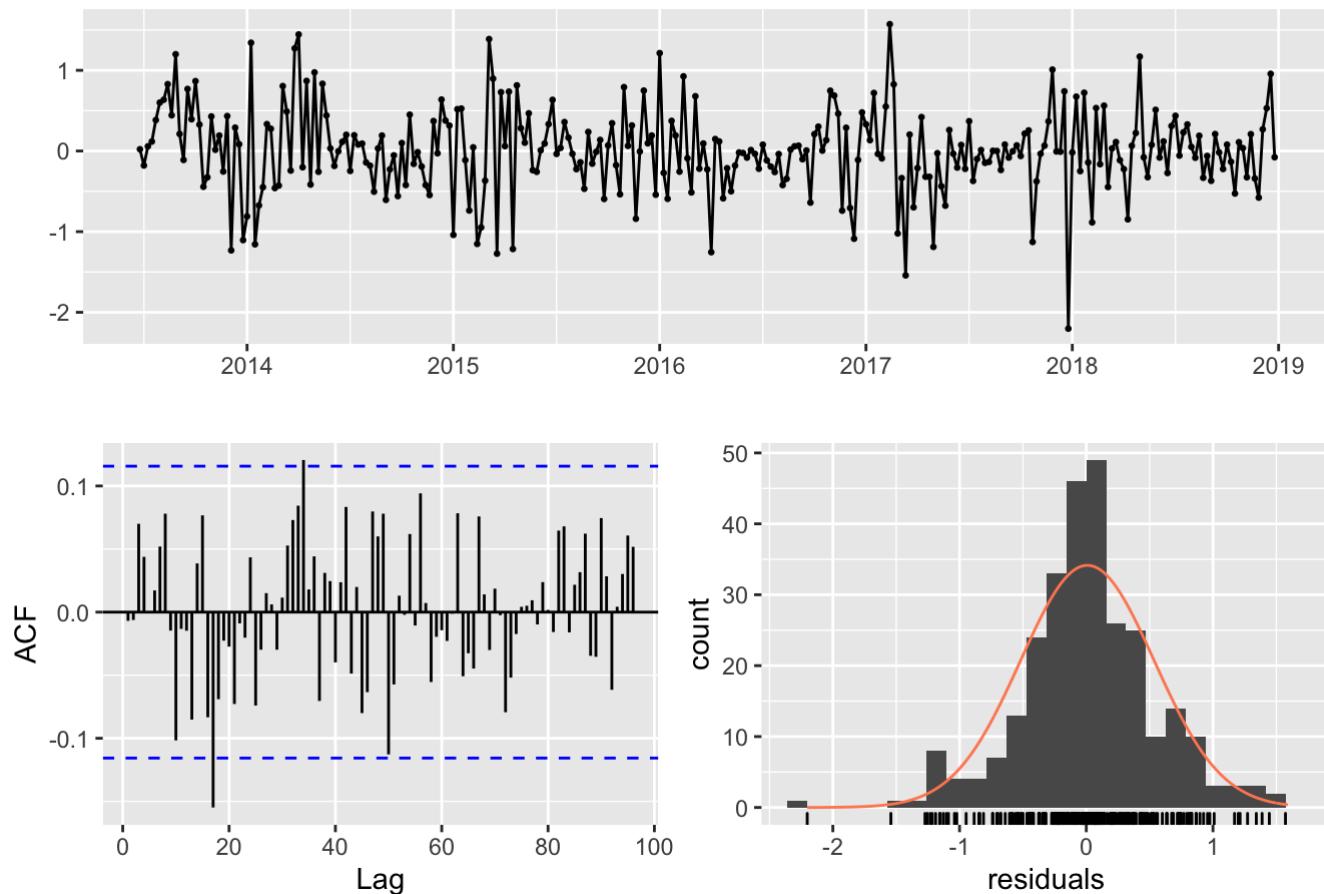


```
summary(arima_fourier_9)
```

```
## Series: train_divvy
## Regression with ARIMA(0,1,2) errors
## Box Cox transformation: lambda= 0.03488702
##
## Coefficients:
##          ma1      ma2     drift    S1-52    C1-52    S2-52    C2-52    S3-52
##        -0.5985  -0.2310  0.0108  1.1181  2.0605 -0.2459 -0.2880 -0.0208
##  s.e.    0.0599   0.0694  0.0055  0.0838  0.0812  0.0616  0.0604  0.0559
##          C3-52    S4-52    C4-52    S5-52    C5-52    S6-52    C6-52    S7-52
##        0.1287  -0.0204  -0.0454 -0.0021  0.018   -0.0939 -0.0214  0.0501
##  s.e.    0.0555   0.0537  0.0533  0.0521  0.052   0.0510  0.0509  0.0498
##          C7-52    S8-52    C8-52    S9-52    C9-52
##        0.0226  -0.1460  -0.0602  0.2042 -0.0413
##  s.e.    0.0499   0.0488  0.0488  0.0476  0.0476
##
## sigma^2 estimated as 0.3002:  log likelihood=-223.36
## AIC=490.73  AICc=494.57  BIC=571.16
##
## Training set error measures:
##           ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
## Training set 640186.7 13605937 9468819 -4.139442 21.95883 0.494634 0.06714366
```

```
checkresiduals(arima_fourier_9)
```

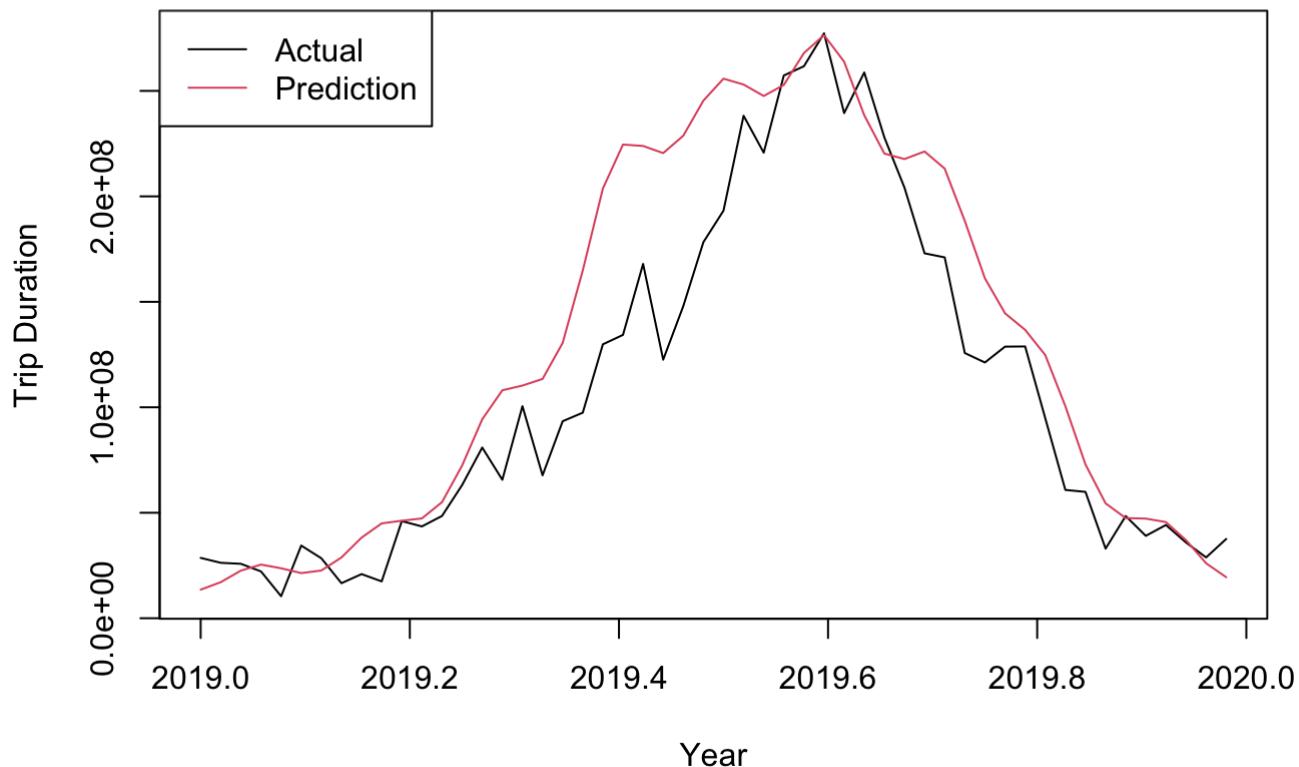
Residuals from Regression with ARIMA(0,1,2) errors



```
##  
## Ljung-Box test  
##  
## data: Residuals from Regression with ARIMA(0,1,2) errors  
## Q* = 64.7, df = 36, p-value = 0.002325  
##  
## Model df: 21. Total lags used: 57
```

We found that when K=9, the AICc is the smallest. (AICc=494.57). However, as the K gets larger, the auto-correlation in the residual becomes more significant. When K=9, the Ljung Box test indicates that there is auto-correlation in its residual. We should try some other parameter.

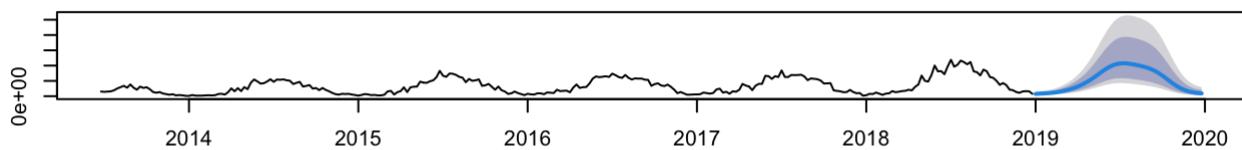
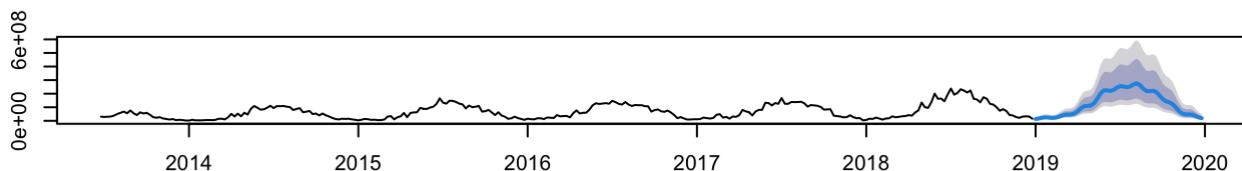
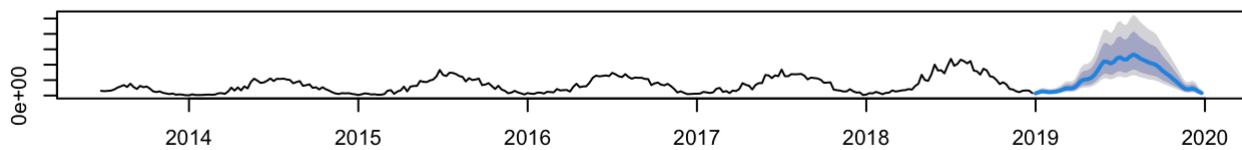
```
plot(test_divvy,col=1,xlab="Year", ylab="Trip Duration")  
lines(arima_fourier_forecast_9$mean,col=2)  
legend('topleft',legend =c('Actual','Prediction'),col=1:2,lty=1)
```



```
MAPE(y_pred=arima_fourier_forecast_9$mean,y_true=test_divvy)
```

```
## [1] 0.3375399
```

```
par(mfrow=c(3,1))
plot(arima_fourier_forecast_3, main='K = 3')
plot(arima_fourier_forecast_9, main='K = 9')
plot(arima_fourier_forecast_13, main='K = 13')
```

K = 3**K = 9****K = 13**

Part 8: TBATS

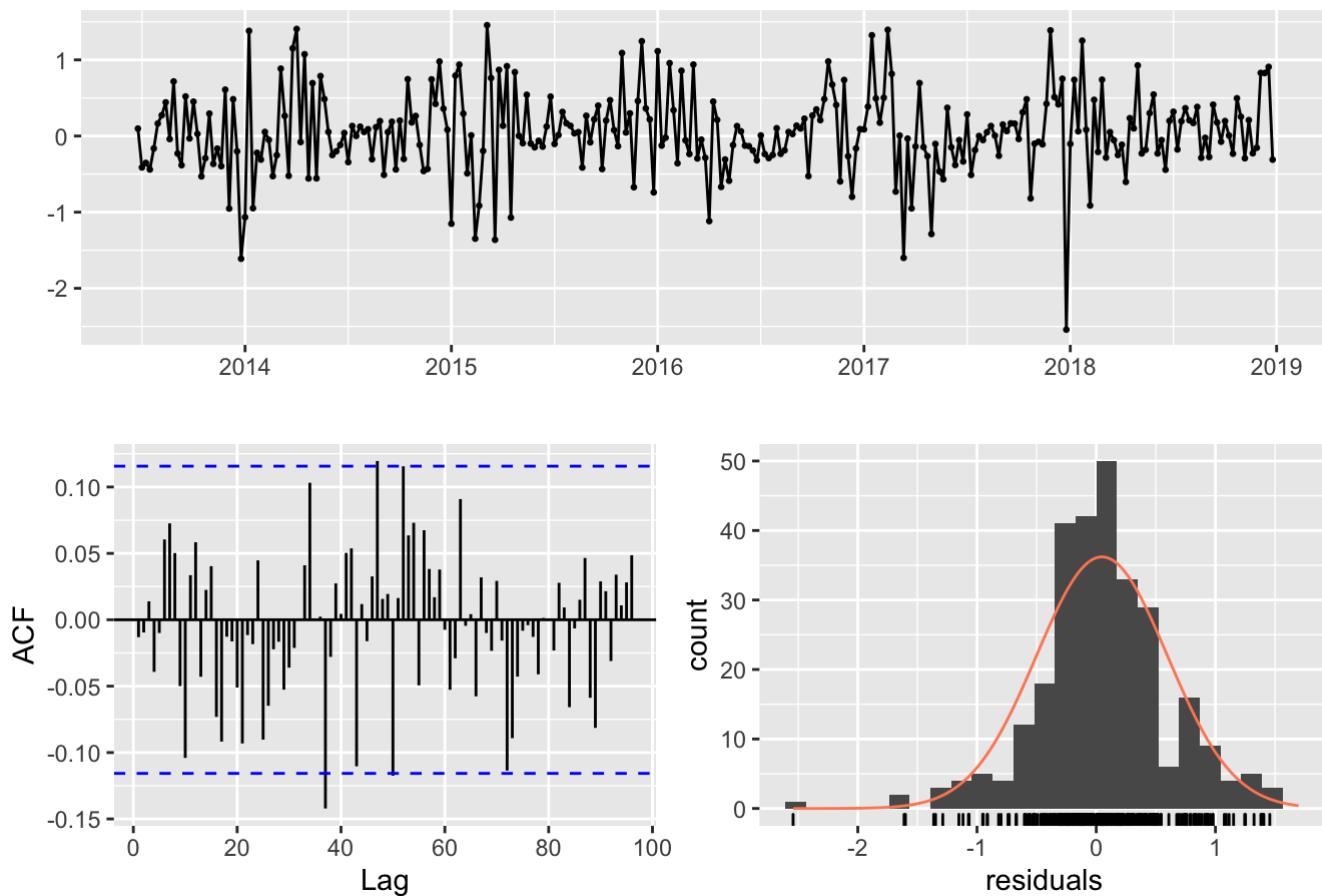
```
tbats_model <- tbats(train_divvy_boxcox, seasonal.period=52)
tbats_model
```

```
## TBATS(1, {2,1}, 0.873, {<52,3>})
##
## Call: tbats(y = train_divvy_boxcox, seasonal.periods = 52)
##
## Parameters
##   Alpha: 0.4153293
##   Beta: -0.03432027
##   Damping Parameter: 0.873104
##   Gamma-1 Values: 0.000541631
##   Gamma-2 Values: -0.0006066662
##   AR coefficients: 0.596125 -0.288484
##   MA coefficients: -0.543414
##
## Seed States:
##           [,1]
## [1,] 20.93749400
## [2,] 0.39032850
## [3,] 2.41049684
## [4,] -0.30060597
## [5,] 0.14151109
## [6,] 0.91360756
## [7,] -0.15710719
## [8,] -0.04656257
## [9,] 0.00000000
## [10,] 0.00000000
## [11,] 0.00000000
##
## Sigma: 0.5509299
## AIC: 1320.083
```

Check residual

```
checkresiduals(tbats_model)
```

Residuals from TBATS

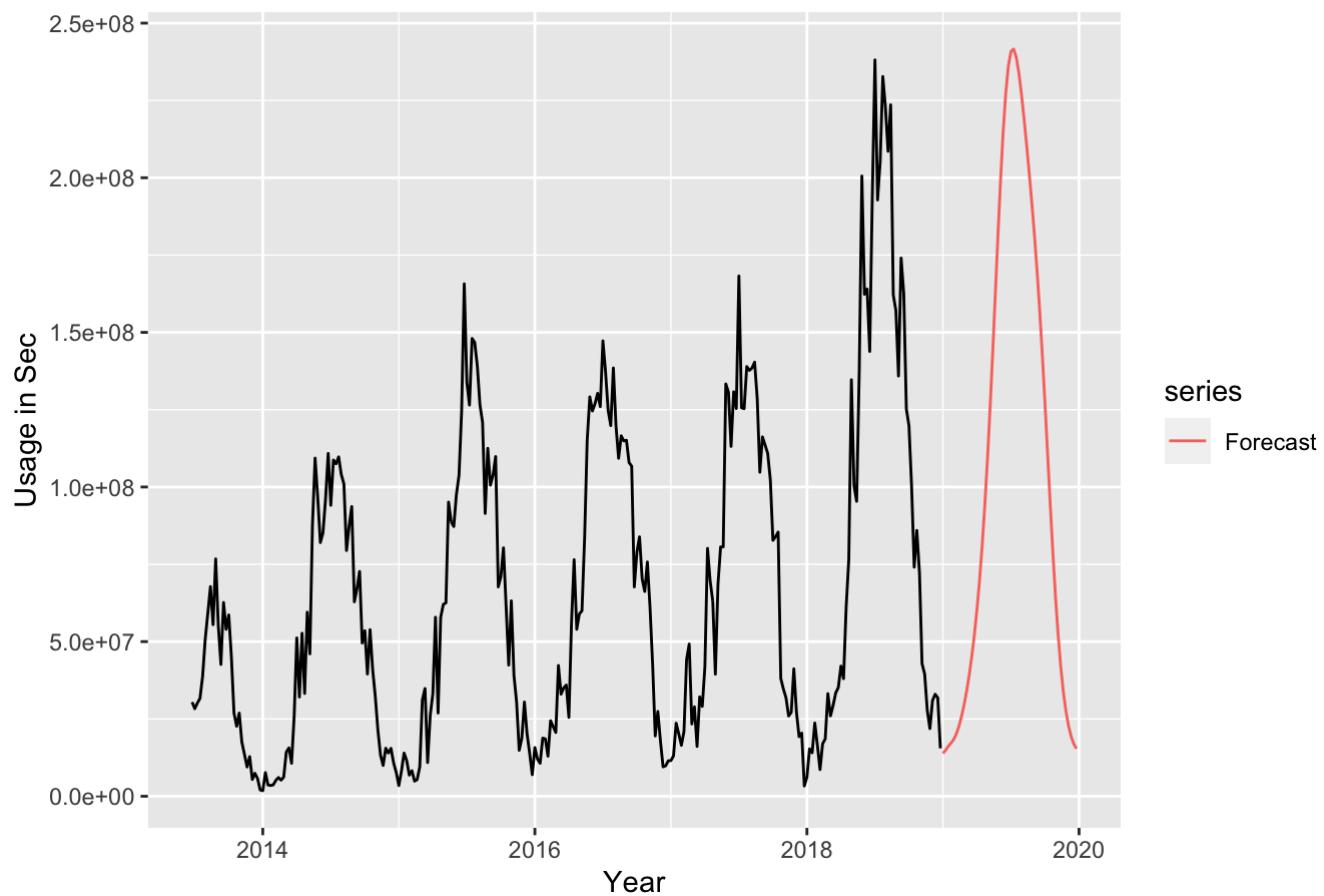


```
##  
## Ljung-Box test  
##  
## data: Residuals from TBATS  
## Q* = 63.714, df = 38, p-value = 0.005579  
##  
## Model df: 19. Total lags used: 57
```

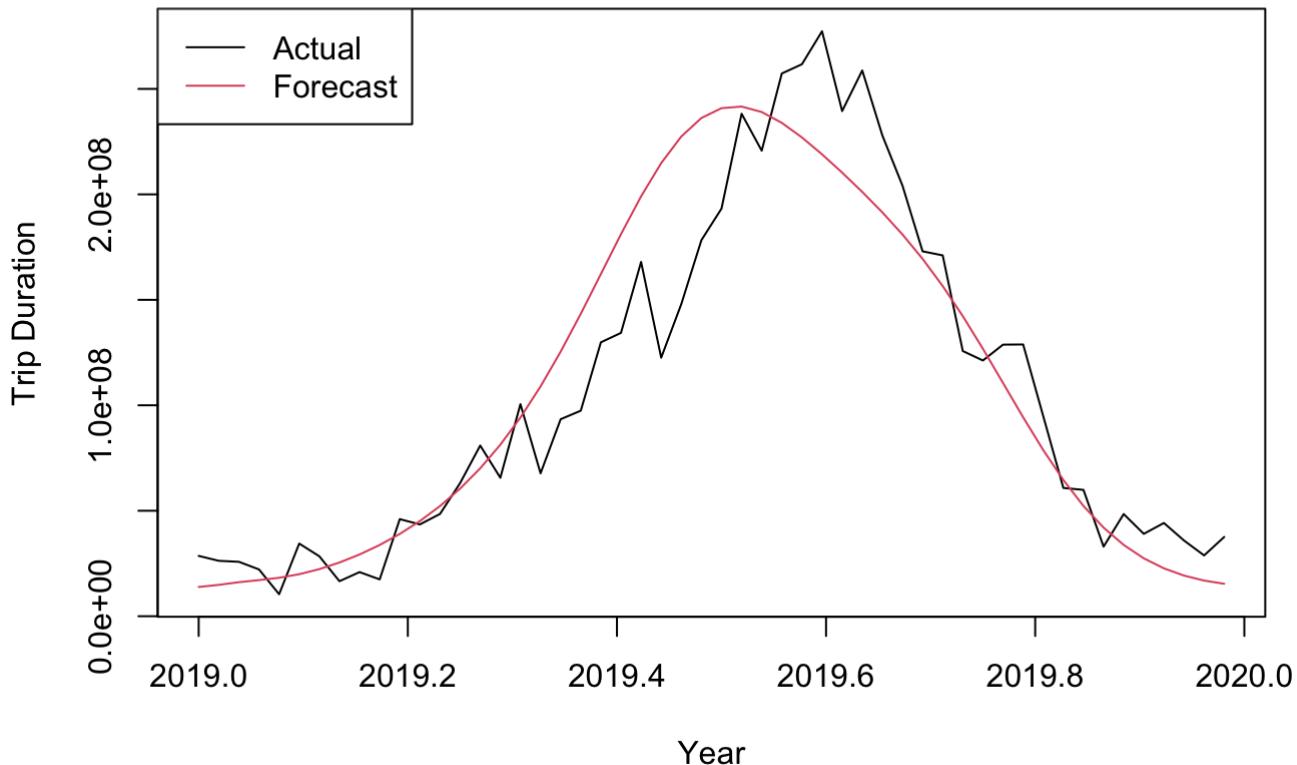
Reverse the BoxCox transformation

```
tbats_model_forecasts <- forecast(tbats_model,h=52)  
forecasts_invboxcox <- InvBoxCox(tbats_model_forecasts$mean,lambda=0.03489689)
```

```
autoplot(train_divvy) +  
  autolayer((forecasts_invboxcox), series="Forecast") +  
  xlab("Year") + ylab("Usage in Sec")
```



```
plot(test_divvy,col=1,xlab="Year", ylab="Trip Duration")
lines(forecasts_invboxcox,col=2)
legend('topleft',legend =c('Actual','Forecast'),col=1:2,lty=1)
```



```
MAPE(y_pred=forecasts_invboxcox,y_true=test_divvy)
```

```
## [1] 0.2815186
```

Part 9: Neural Network Autoregression

```
nnetar_model <- nnetar(train_divvy,lambda=0)
nnetar_model
```

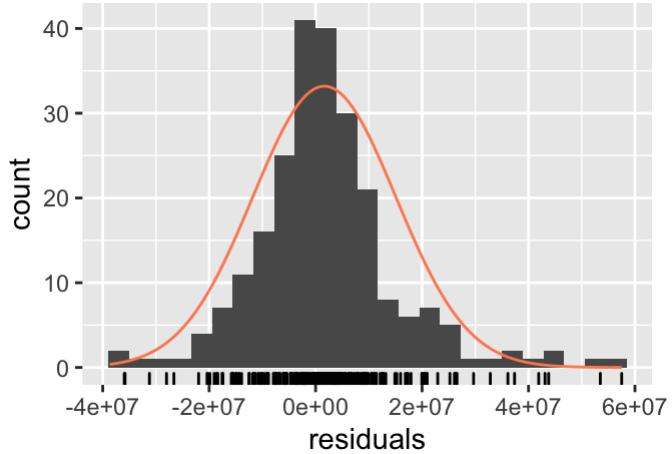
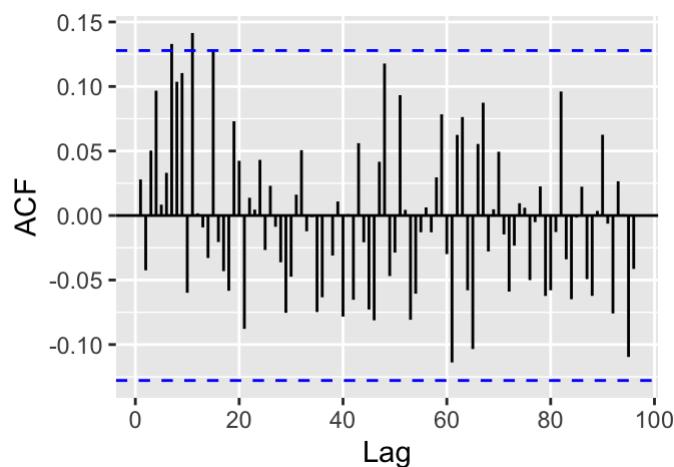
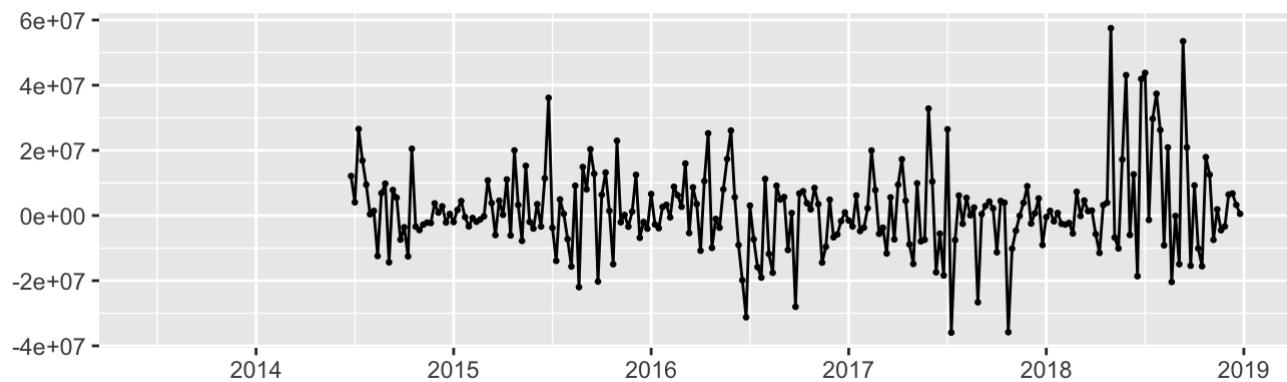
```
## Series: train_divvy
## Model: NNDAR(7,1,4)[52]
## Call: nnetar(y = train_divvy, lambda = 0)
##
## Average of 20 networks, each of which is
## a 8-4-1 network with 41 weights
## options were - linear output units
##
## sigma^2 estimated as 0.06479
```

Check residual

```
checkresiduals(nnetar_model$residuals)
```

```
## Warning in modeldf.default(object): Could not find appropriate degrees of
## freedom for this model.
```

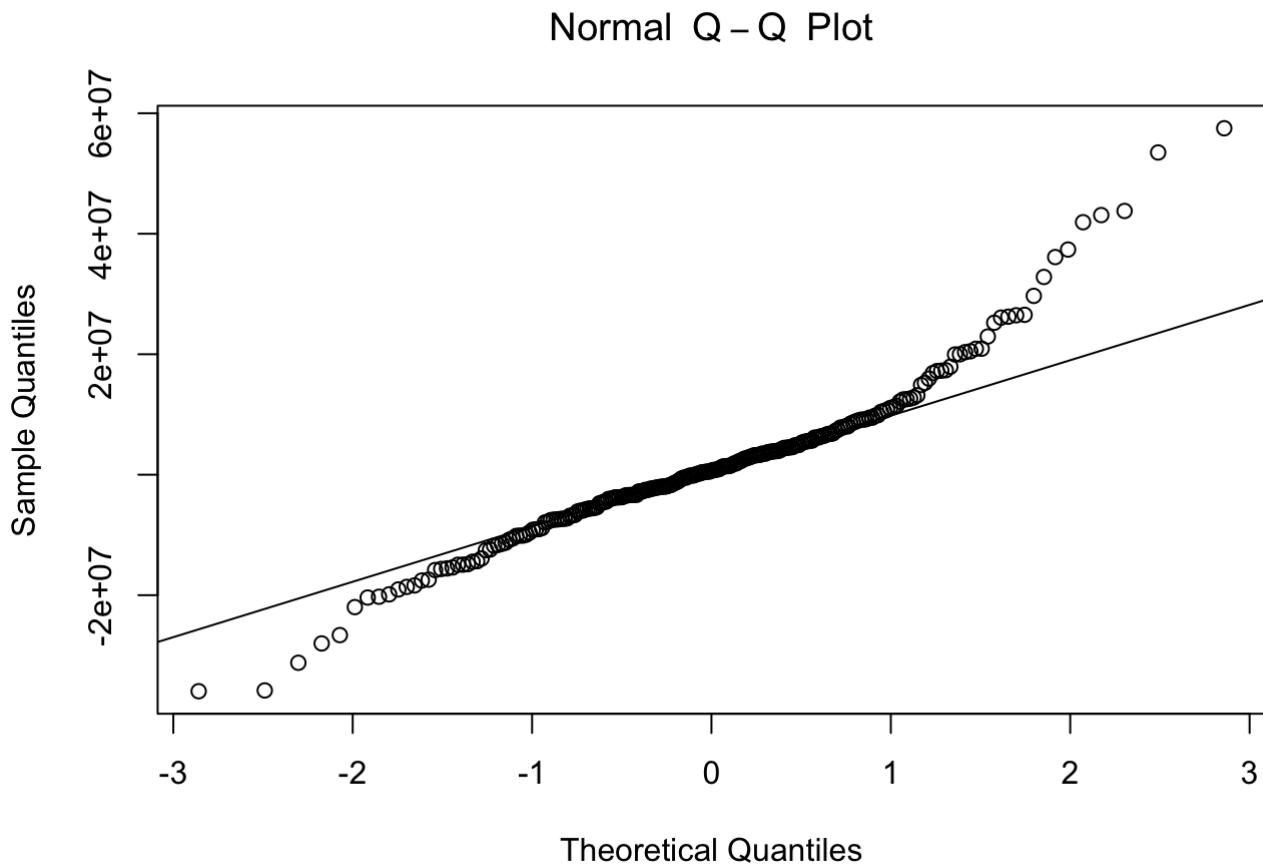
Residuals



```
shapiro.test(nnetar_model$residuals)
```

```
##
## Shapiro-Wilk normality test
##
## data: nnetar_model$residuals
## W = 0.94199, p-value = 4.843e-08
```

```
qqnorm(nnetar_model$residuals, main=expression(Normal~~Q-Q~~Plot))
qqline(nnetar_model$residuals)
```



```
Box.test(nnetar_model$residuals,lag=52,type = c("Ljung-Box"),fitdf=41)
```

```
##  
## Box-Ljung test  
##  
## data: nnetar_model$residuals  
## X-squared = 52.185, df = 11, p-value = 2.524e-07
```

Part 10: ETS

Build model

```
ets_fit <- stlf(train_divvy)  
summary(ets_fit)
```

```

##  

## Forecast method: STL + ETS(A,Ad,N)  

##  

## Model Information:  

## ETS(A,Ad,N)  

##  

## Call:  

## ets(y = na.interp(x), model = etsmodel, allow.multiplicative.trend = allow.multiplic  

##     ative.trend)  

##  

## Smoothing parameters:  

##   alpha = 0.4015  

##   beta  = 1e-04  

##   phi   = 0.9346  

##  

## Initial states:  

##   l = -45225995.0126  

##   b = 7244827.9108  

##  

## sigma: 12469859  

##  

##      AIC      AICc      BIC  

## 11009.71 11010.01 11031.67  

##  

## Error measures:  

##  

##          ME      RMSE      MAE      MPE      MAPE      MASE      ACF1  

## Training set 213037.2 12360759 9001462 -8.124831 25.64128 0.4702201 0.09706505  

##  

## Forecasts:  

##  

##       Point Forecast    Lo 80     Hi 80      Lo 95     Hi 95  

## 2019.000 27998085 12017318 43978852 3557610.836 52438559  

## 2019.019 30955263 13734177 48176349 4617884.953 57292641  

## 2019.038 32842445 14464071 51220819 4735147.484 60949743  

## 2019.058 35062086 15594671 54529501 5289243.615 64834928  

## 2019.077 31285270 10786208 51784332 -65340.053 62635879  

## 2019.096 29390426 7908864 50871987 -3462787.311 62243639  

## 2019.115 37964849 15543495 60386204 3674347.071 72255352  

## 2019.135 38570510 15246925 61894094 2900165.730 74240854  

## 2019.154 36209308 12016871 60401745 -789831.658 73208447  

## 2019.173 45461226 20429854 70492599 7179045.641 83743407  

## 2019.192 41975182 16131891 67818473 2451279.372 81499084  

## 2019.212 41287767 14657117 67918417 559702.540 82015831  

## 2019.231 46088209 18692653 73483765 4190322.148 87986097  

## 2019.250 53109055 24969226 81248885 10072899.878 96145210  

## 2019.269 68791484 39926428 97656541 24646189.982 112936779  

## 2019.288 72891389 43318759 102464019 27663954.763 118118824  

## 2019.308 72723330 42459550 102987109 26438872.756 119007787  

## 2019.327 85991637 55052038 116931236 38673603.171 133309671  

## 2019.346 82817938 51216873 114419004 34488279.744 131147597  

## 2019.365 101640170 69391115 133889225 52319497.012 150960844  

## 2019.385 119667169 86782809 152551528 69374881.073 169959456  

## 2019.404 146102022 112594327 179609716 94856425.272 197347618  

## 2019.423 135550758 101431046 169670470 83369161.215 187732354

```

## 2019.442	134009656	99288650	168730662	80908459.619	187110852
## 2019.462	139528969	104216849	174841089	85523742.036	193534196
## 2019.481	153424477	117530924	189318031	98530025.148	208318930
## 2019.500	166917716	130451950	203383481	111148140.229	222687291
## 2019.519	147715070	110685888	184744251	91083823.941	204346315
## 2019.538	151186023	113601830	188770216	93705959.891	208666086
## 2019.558	159049909	120918744	197181074	100733324.276	217366494
## 2019.577	159102085	120431648	197772521	99960755.301	218243414
## 2019.596	150615587	111413264	189817910	90660808.413	210570366
## 2019.615	148596270	108869152	188323388	87838886.556	209353654
## 2019.635	129721545	89476447	169966642	68171978.872	191271110
## 2019.654	129636678	88880157	170393198	67304958.458	191968397
## 2019.673	119205301	77943672	160466929	56101085.939	182309515
## 2019.692	124816827	83056178	166577475	60949426.355	188684227
## 2019.712	124991557	82737761	167245352	60369952.695	189613160
## 2019.731	96507326	53766054	139248597	31140191.819	161874460
## 2019.750	94831719	51608452	138054986	28727436.075	160936001
## 2019.769	90267655	46567693	133967618	23434330.144	157100981
## 2019.788	79375354	35203826	123546883	11820831.420	146929877
## 2019.808	65681138	21043012	110319264	-2586984.912	133949260
## 2019.827	66520519	21420610	111620428	-2453839.851	135494878
## 2019.846	50181839	4624816	95738863	-19491615.399	119855294
## 2019.865	41424581	-4585025	87434188	-28941040.053	111790203
## 2019.885	30821250	-15636542	77279041	-40229810.643	101872310
## 2019.904	35359201	-11542501	82260904	-36370762.646	107089165
## 2019.923	33313116	-14028344	80654575	-39089398.766	105715630
## 2019.942	28752696	-19024481	76529873	-44316191.263	101821583
## 2019.962	26549023	-21659941	74757987	-47180225.425	100278271
## 2019.981	17744512	-30892412	66381437	-56639244.858	92128269
## 2020.000	27948736	-21112422	77009895	-47083830.313	102981303
## 2020.019	30909144	-18572617	80390905	-44766678.489	106584967
## 2020.038	32799344	-17099479	82698168	-43514320.563	109113009
## 2020.058	35021805	-15290629	85334239	-41924422.274	111968033
## 2020.077	31247625	-19475051	81970302	-46326013.951	108821264
## 2020.096	29355245	-21774387	80484876	-48840778.504	107551268
## 2020.115	37931971	-13601406	89465348	-40881528.268	116745470
## 2020.135	38539783	-13394205	90473770	-40886397.340	117965962
## 2020.154	36180591	-16150943	88512126	-43853584.481	116214767
## 2020.173	45434389	-7291699	98160477	-35203203.375	126071982
## 2020.192	41950101	-11167612	95067814	-39286430.902	123186633
## 2020.212	41264327	-12242147	94770802	-40566764.206	123095419
## 2020.231	46066304	-7826131	99958738	-36355062.520	128487670
## 2020.250	53088583	-1187069	107364235	-29918864.181	136096030
## 2020.269	68772352	14116167	123428537	-14817070.174	152361774
## 2020.288	72873509	17839420	127907598	-11293867.423	157040885
## 2020.308	72706619	17297202	128116037	-12034772.266	157448011
## 2020.327	85976020	30193798	141758243	664472.601	171287568
## 2020.346	82803344	26650790	138955898	-3074576.831	168681264
## 2020.365	101626531	45106070	158146991	15185945.802	188067116
## 2020.385	119654421	62768432	176540410	32654808.920	206654034
## 2020.404	146090109	88840924	203339294	58535036.011	233645182
## 2020.423	135539624	77929532	193149717	47432591.135	223646658
## 2020.442	133999251	76030497	191968005	45343691.951	222654810
## 2020.462	139519245	81194034	197844456	50318531.753	228719958

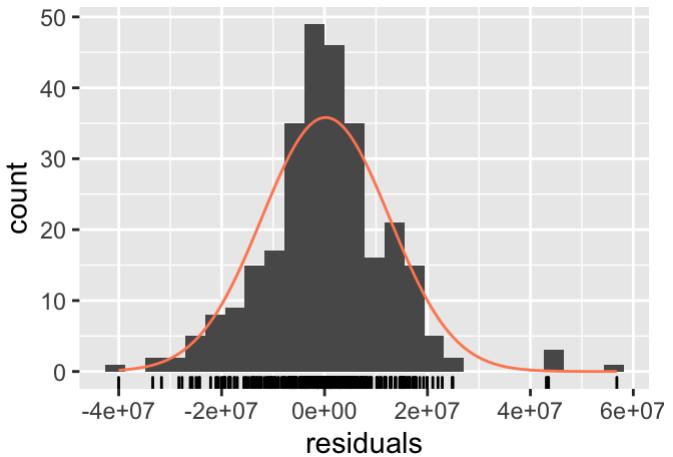
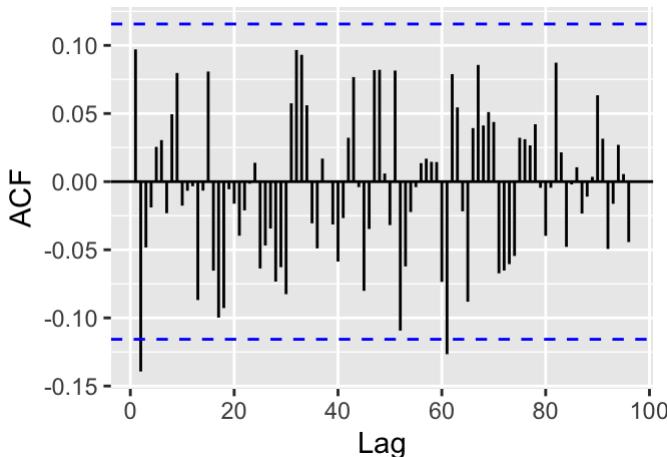
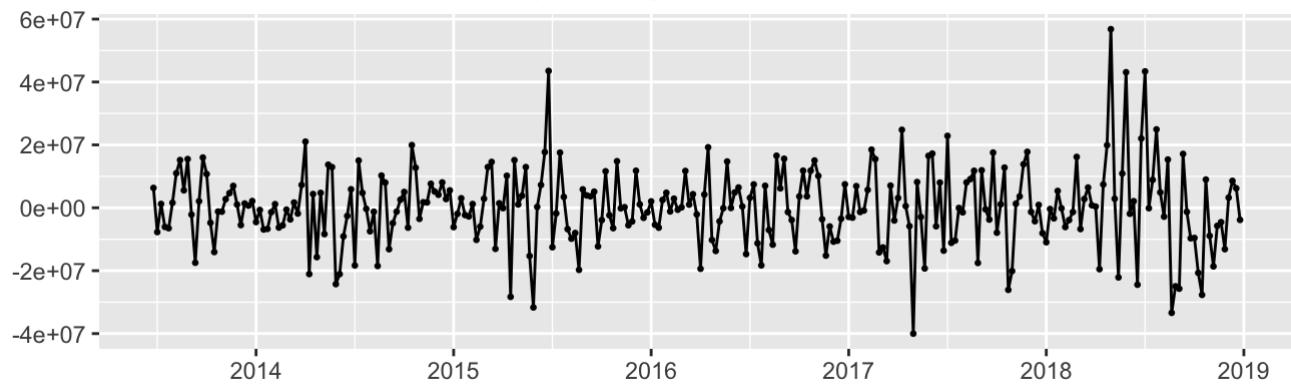
```
## 2020.481    153415390  94735886 212094894  63672832.540 243157947
## 2020.500    166909223  107877551 225940895  76628071.993 257190374
## 2020.519    147707132  88325381 207088884  56890580.736 238523684
## 2020.538    151178605  91448825 210908385  59829789.484 242527421
## 2020.558    159042977  98967184 219118770  67164979.538 250920974
## 2020.577    159095606  98675780 219515432  66691456.441 251499755
## 2020.596    150609532  89847622 211371443  57682208.914 243536856
## 2020.615    148590612  87488530 209692693  55143042.370 242038181
## 2020.635    129716256  68275887 191156626  35751320.487 223681192
## 2020.654    129631736  67854931 191408541  35152265.503 224111206
## 2020.673    119200682  57089263 181312101  24209464.051 214191900
## 2020.692    124812510  62368270 187256751  29312286.045 220312735
## 2020.712    124987523  62212225 187762820  28980990.016 220994055
## 2020.731    96503556   33398938 159608174  -6629.369 193013741
## 2020.750    94828195   31395966 158260425  -2183028.245 191839419
## 2020.769    90264363   26506205 154022521  -7245325.410 187774051
## 2020.788    79372277   15289848 143454706  -18633340.778 177377895
## 2020.808    65678262   1273194 130083330  -32820789.242 164177313
## 2020.827    66517832   1791733 131243930  -32472193.695 165507857
## 2020.846    50179328  -14866217 115224872  -49299248.983 149657904
## 2020.865    41422234  -23941197 106785664  -58542506.600 141386974
## 2020.885    30819056  -34860722 96498834  -69629495.818 131267608
## 2020.904    35357151  -30637458 101351760  -65572893.336 136287195
## 2020.923    33311200  -32996745 99619145  -68098051.265 134720451
## 2020.942    28750905  -37868903 95370713  -73135298.906 130637109
## 2020.962    26547349  -40382869 93477567  -75813585.715 128908284
## 2020.981    17742948  -49496246 84982143  -85090526.138 120576423
```

Check residuals

```
checkresiduals(ets_fit)
```

```
## Warning in checkresiduals(ets_fit): The fitted degrees of freedom is based on
## the model used for the seasonally adjusted data.
```

Residuals from STL + ETS(A,Ad,N)



```
##  
## Ljung-Box test  
##  
## data: Residuals from STL + ETS(A,Ad,N)  
## Q* = 61.321, df = 52, p-value = 0.1764  
##  
## Model df: 5. Total lags used: 57
```

Calculate MAPE

```
MAPE(ets_fit$mean,test_divvy)
```

```
## [1] 0.3259718
```